

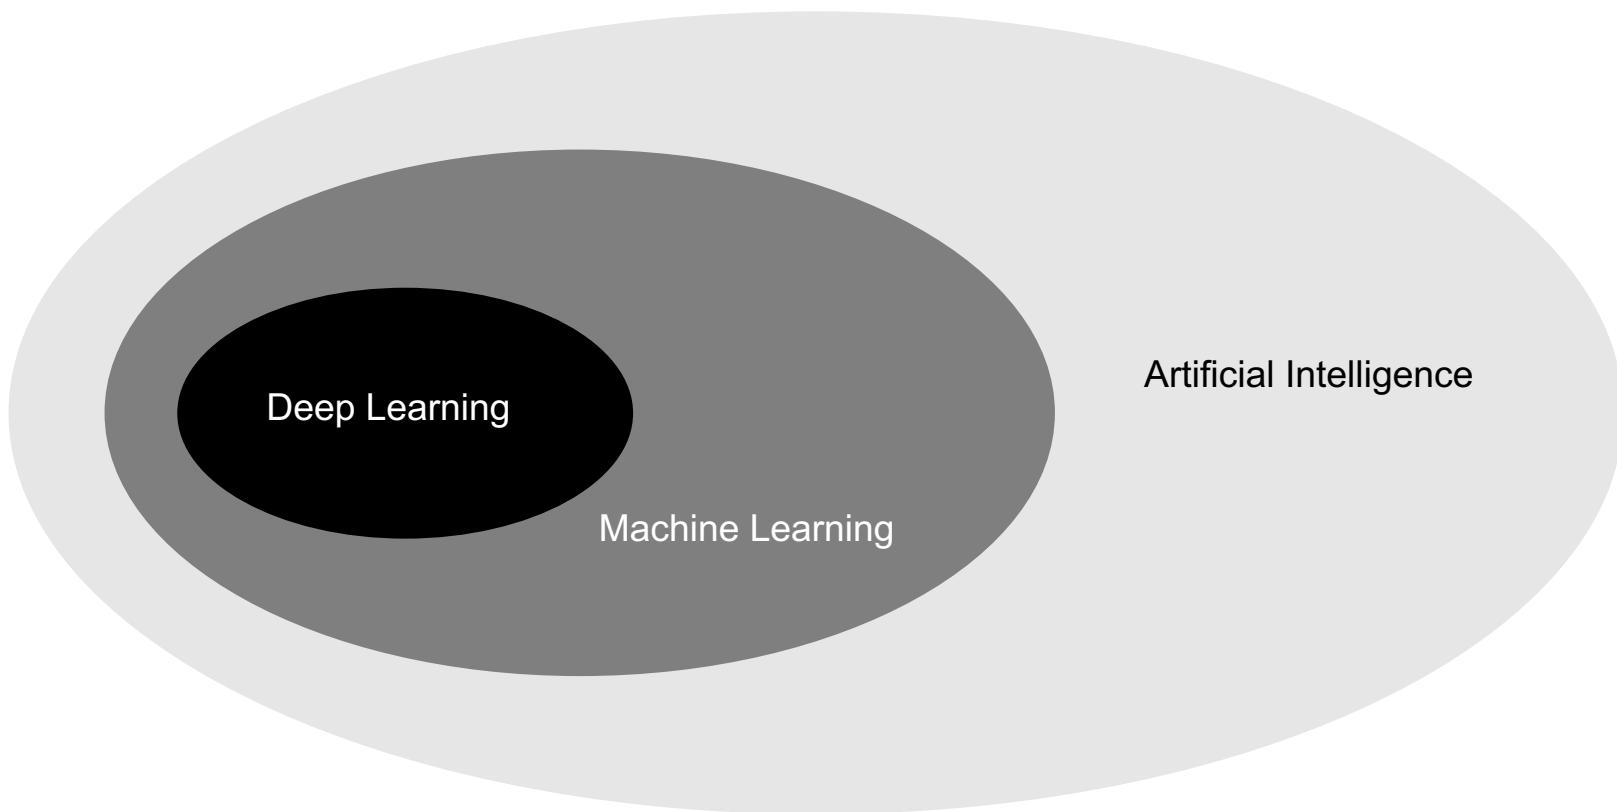
Computational MRI

Machine learning in MRI and neural network
architecture design

Outline

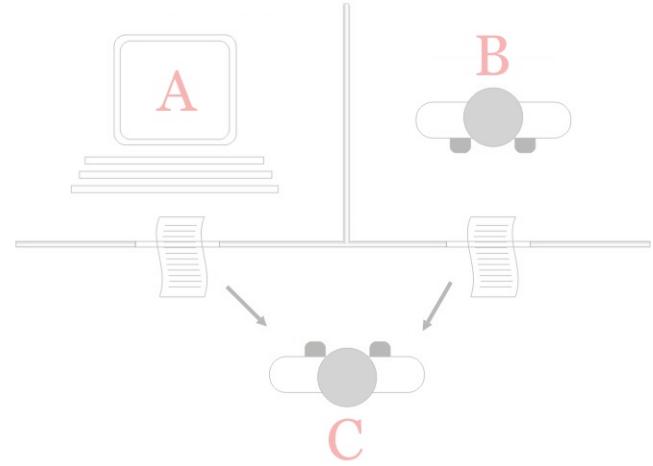
- Overview of ML, DL and AI
- Didactic implementation examples (lab exercise)
 - Linear regression
 - Classification of DTI data (MLP)
 - Classification of reconstructed image quality (CNN)
- Generalization: Model complexity, overfitting vs. underfitting
- Architectures: Convolutional vs fully connected

What is AI?

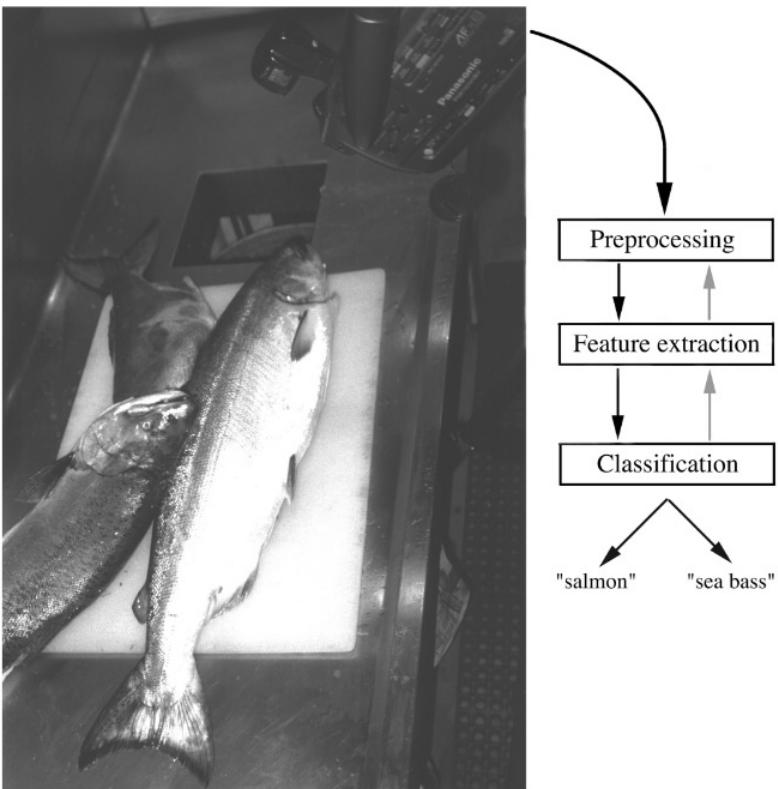


General AI

- Turing Test (Turing)
- Coffee Test (Wozniak)
- Robot College Student Test (Goertzel)
- The Employment Test (Nilsson)
- The flat pack furniture test (Severyns)

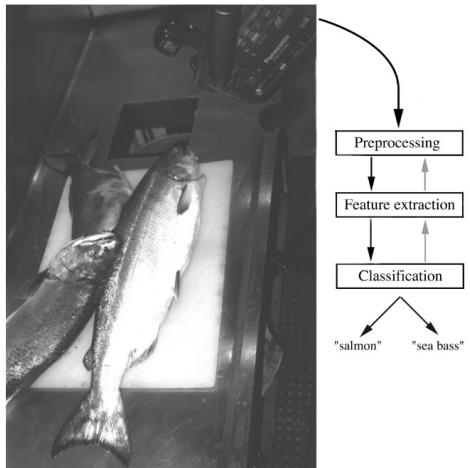


ML: A classic computer vision example

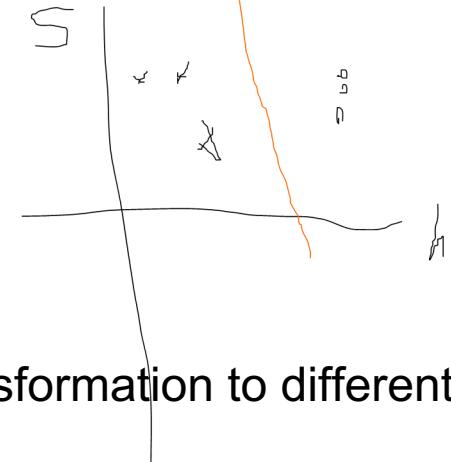


Duda, Hart, Stork: Pattern Classification

Solution 1: Classic computer vision



1. Preprocessing: Find individual fish, rotation, lighting
2. Feature extraction: Fit ellipse
3. Classification:
 - linear classifier: $kx+d$
 - nonlinear classifier or transformation to different feature space



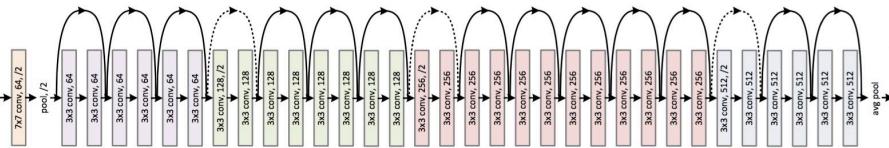
若是非线性DB，
先把feature map 到新的feature space

Solution 2: Deep Learning



34-layer residual

Image



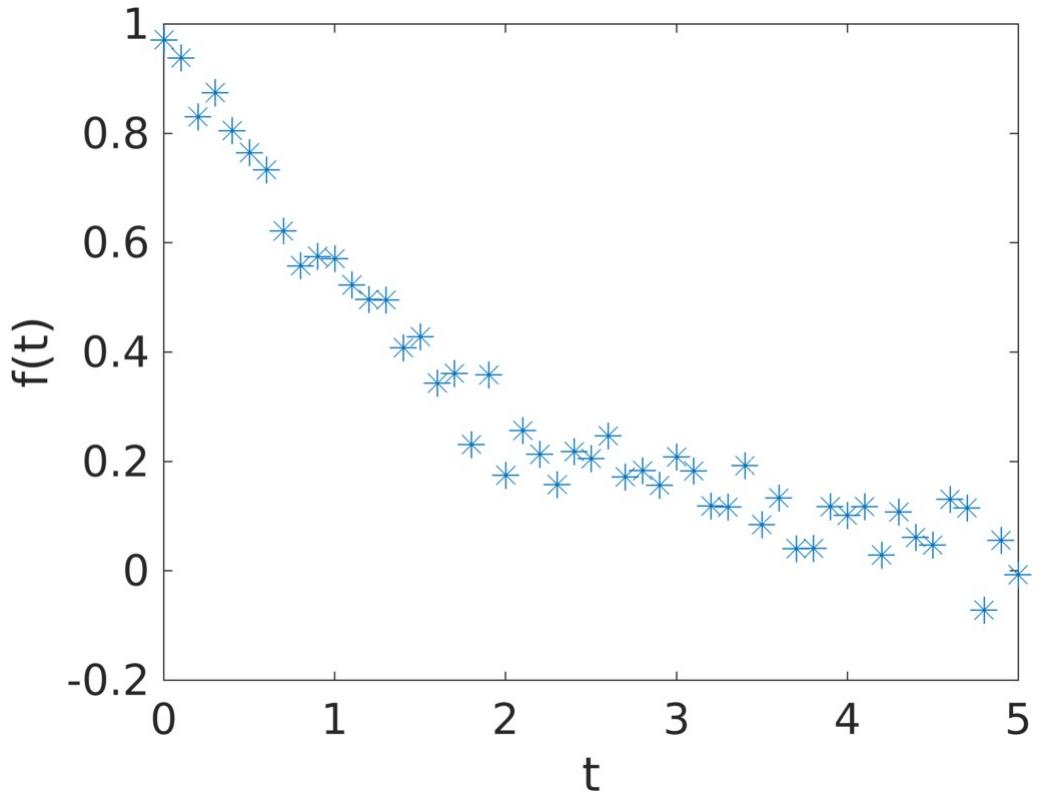
Salmon
Sea bass

Machine Learning 1: Data

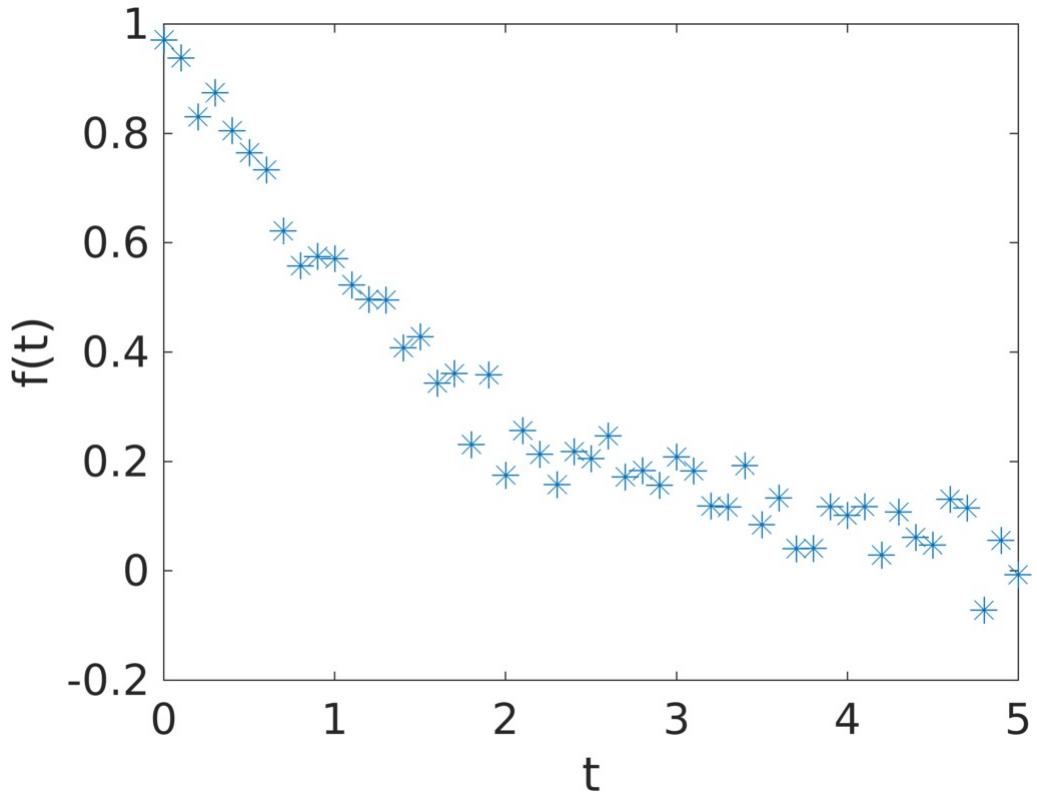


<http://www.image-net.org/>

Fitting a model to data 1/3

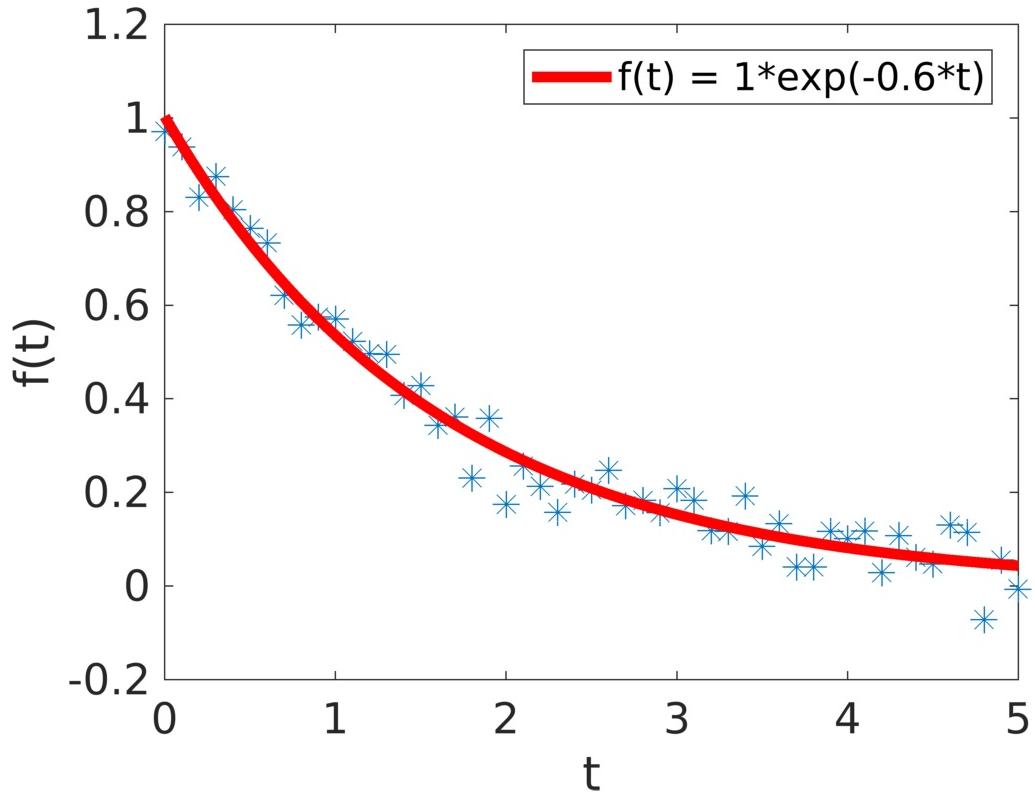


Fitting a model to data 2/3



$$f(t) = ae^{bt}$$

Fitting a model to data 3/3

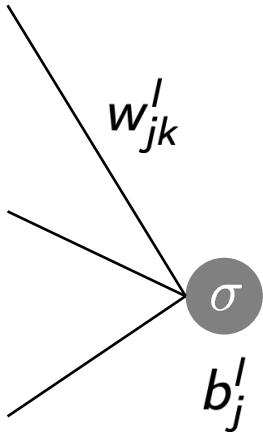


$$f(t) = ae^{bt}$$

$$a = 1$$

$$b = -0.6$$

Neural networks



$$a_j^I = \sigma \left(\sum_k w_{jk}^I a_k^{I-1} + b_j^I \right)$$

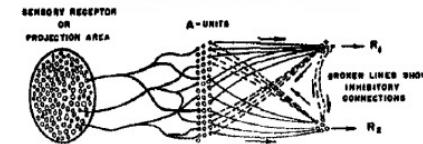


FIG. 2A. Schematic representation of connections in a simple perceptron.

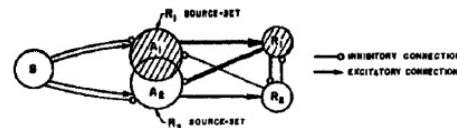
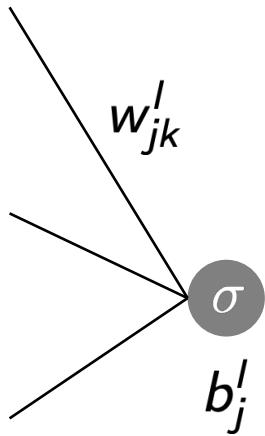


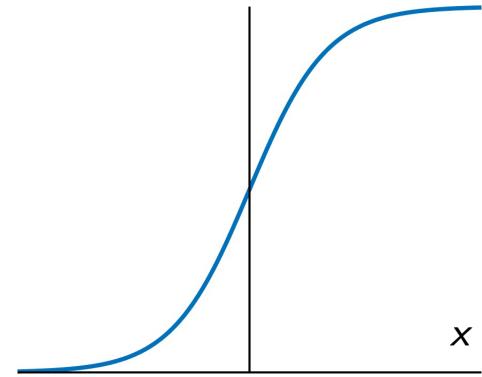
FIG. 2B. Venn diagram of the same perceptron (shading shows active sets for R_1 response).

Neural networks

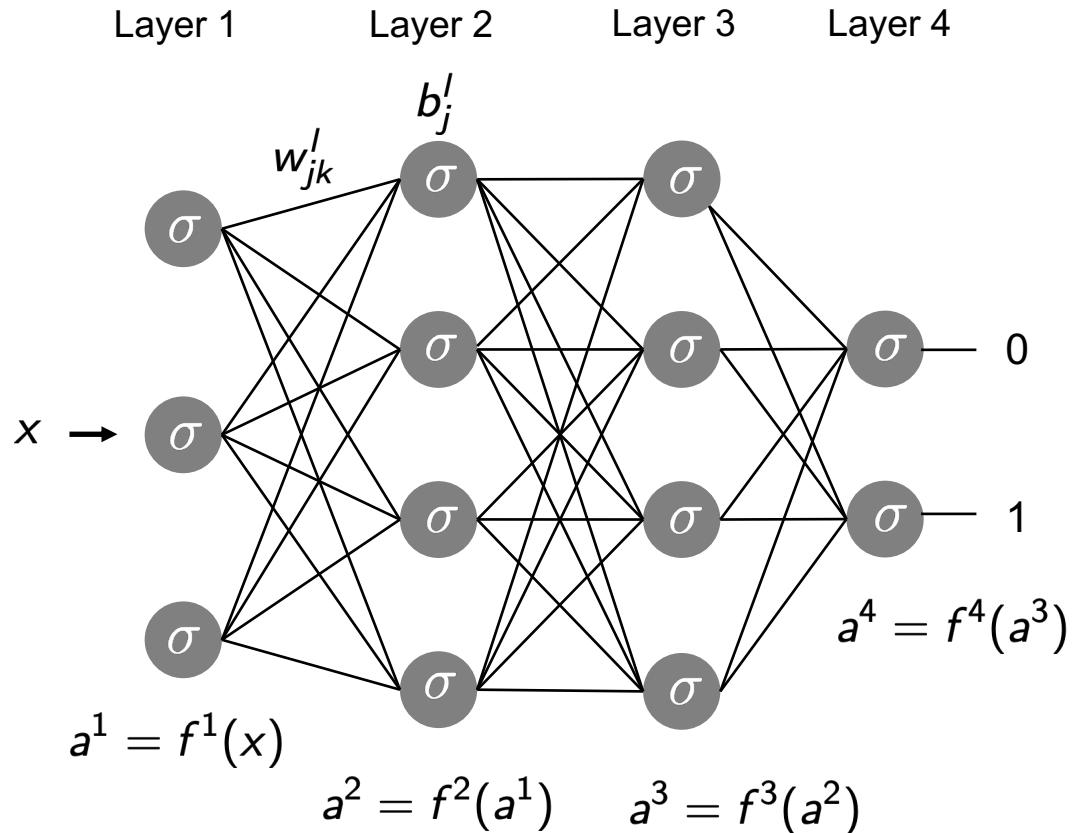


$$a_j^I = \sigma \left(\sum_k w_{jk}^I a_k^{I-1} + b_j^I \right)$$

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



Neural networks



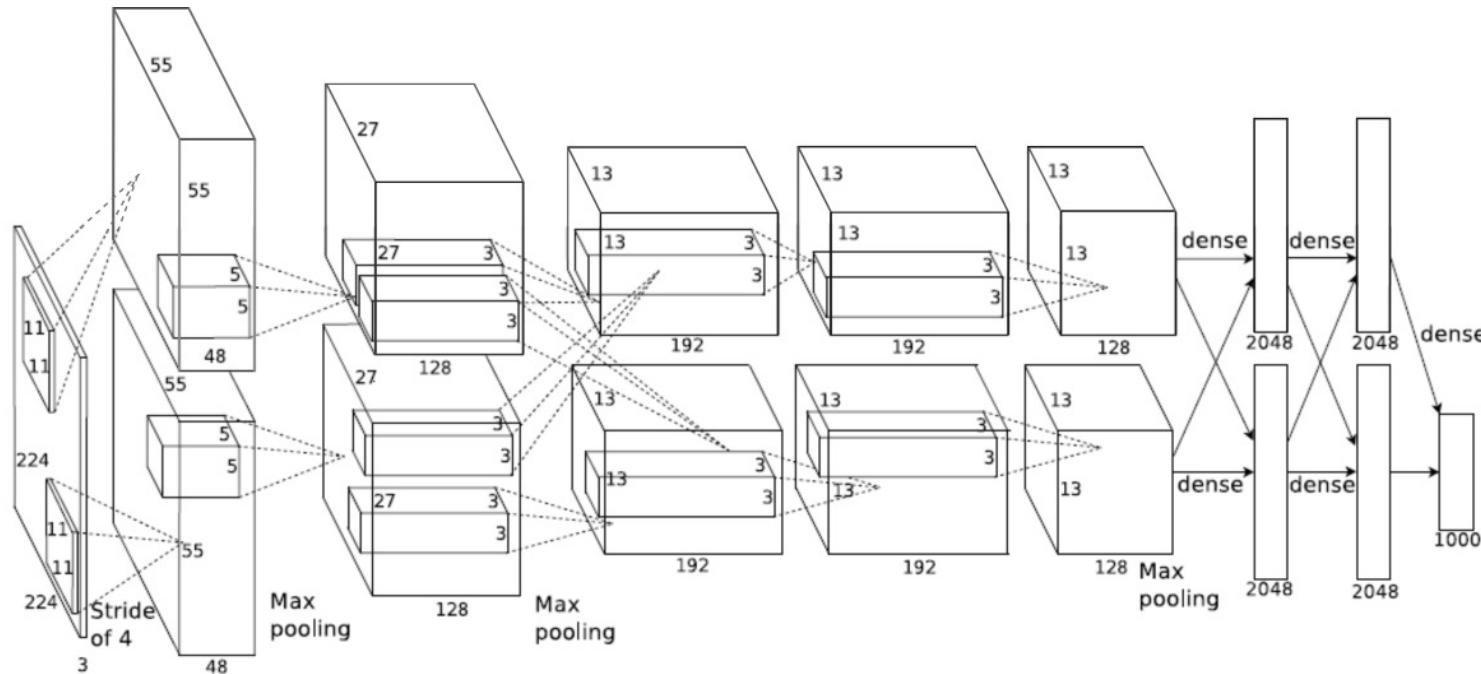
Large number of model parameters

High descriptive capacity

$$a^4 = F(x) = f^4(f^3(f^2(f^1(x))))$$

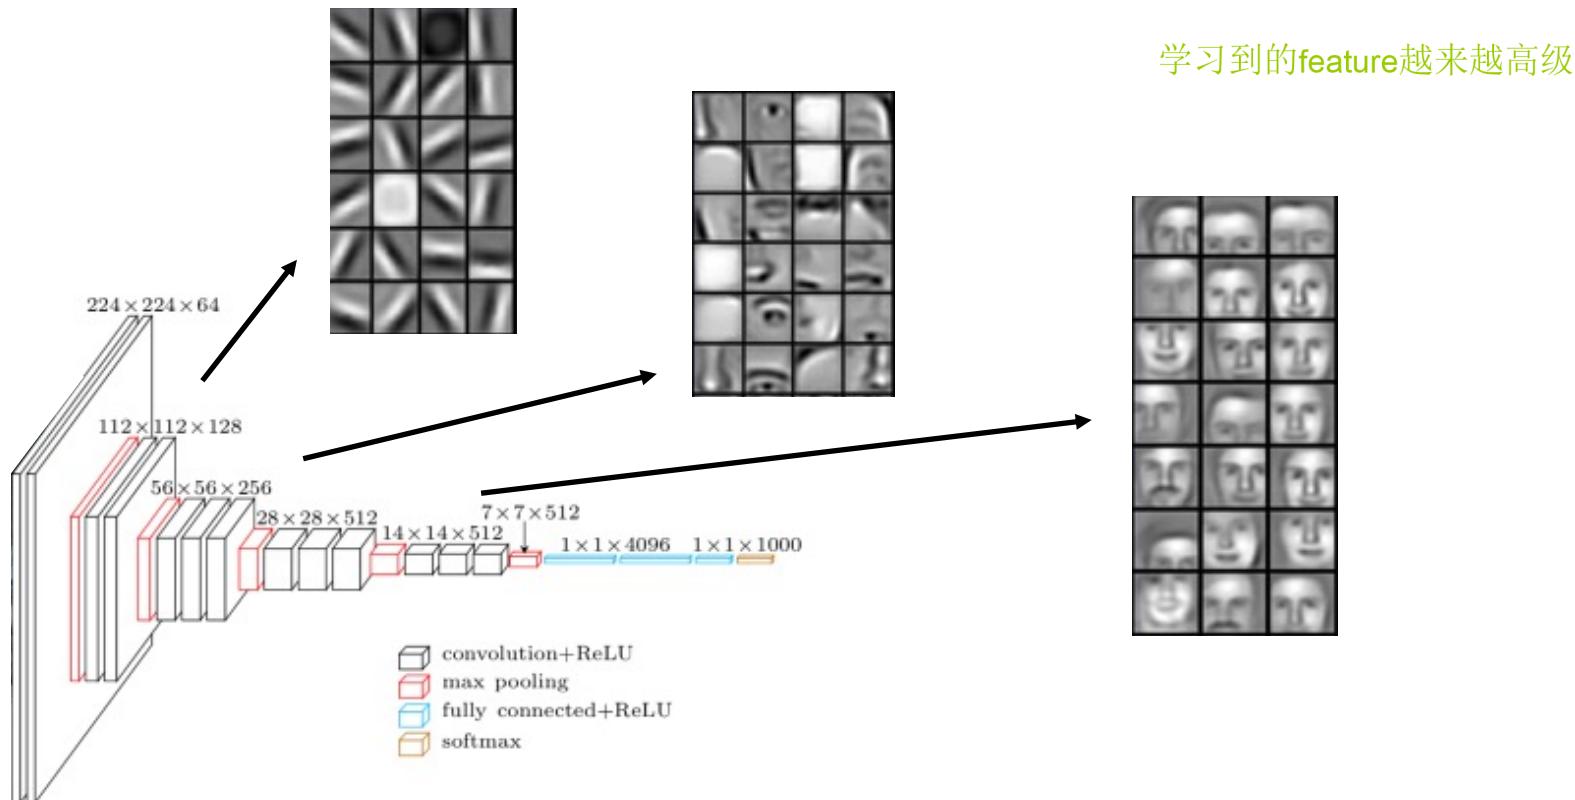
Cybenko: Math Control Sig Sys 1989

Alexnet

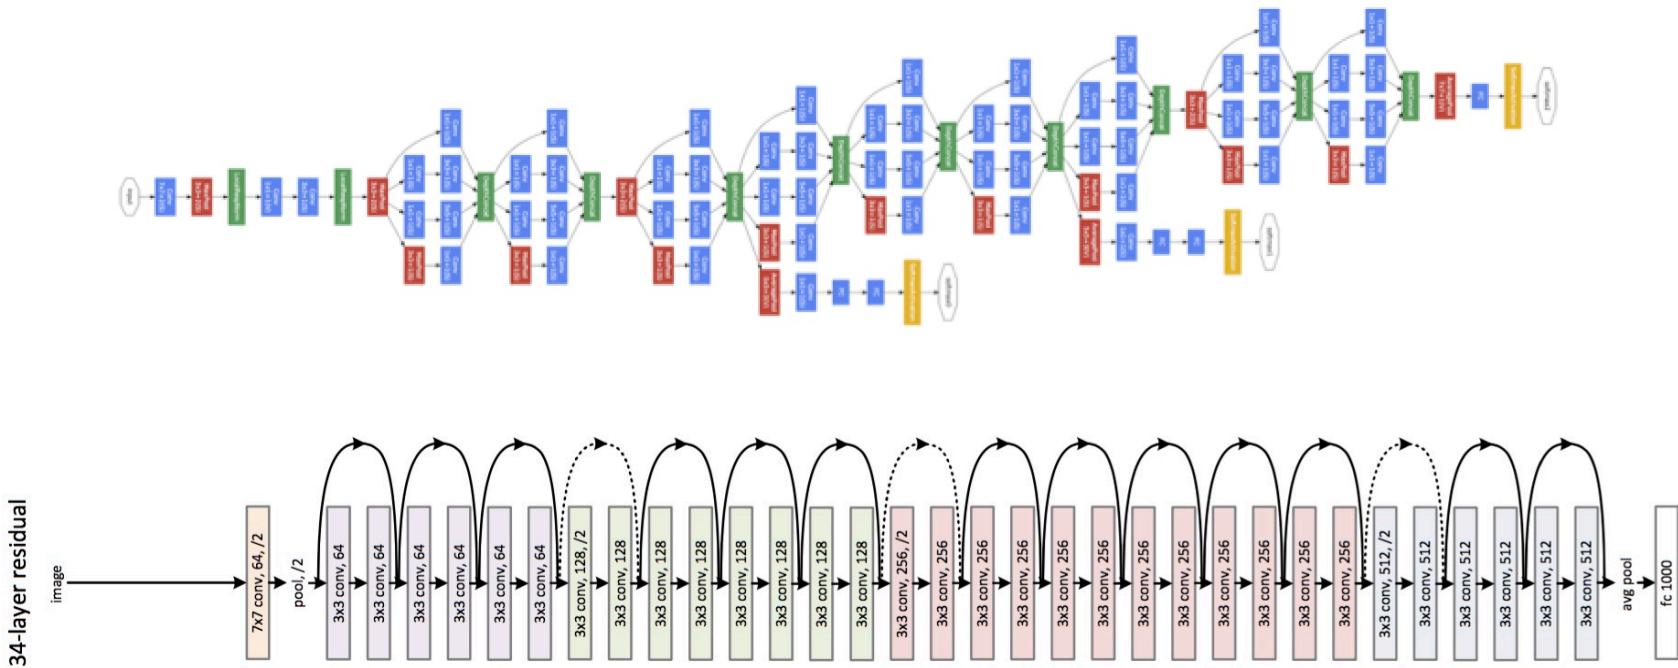


Krizhevsky et al. NIPS 2012

CNN for image categorization



Googlenet (Inception), Resnet

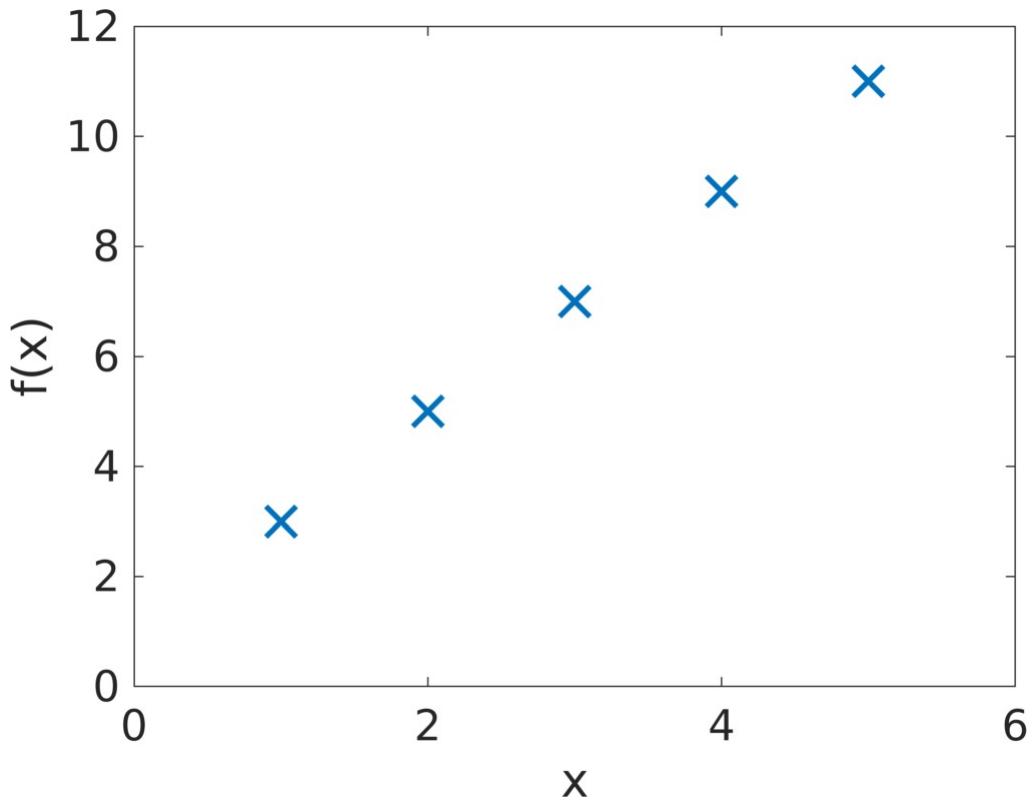


Szegedy et al., CVPR 2015
He et al., CVPR 2016

Example 1: Linear Regression

A low level view of PyTorch

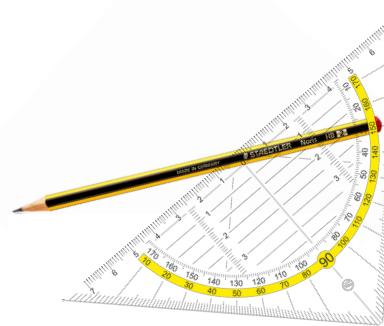
Data



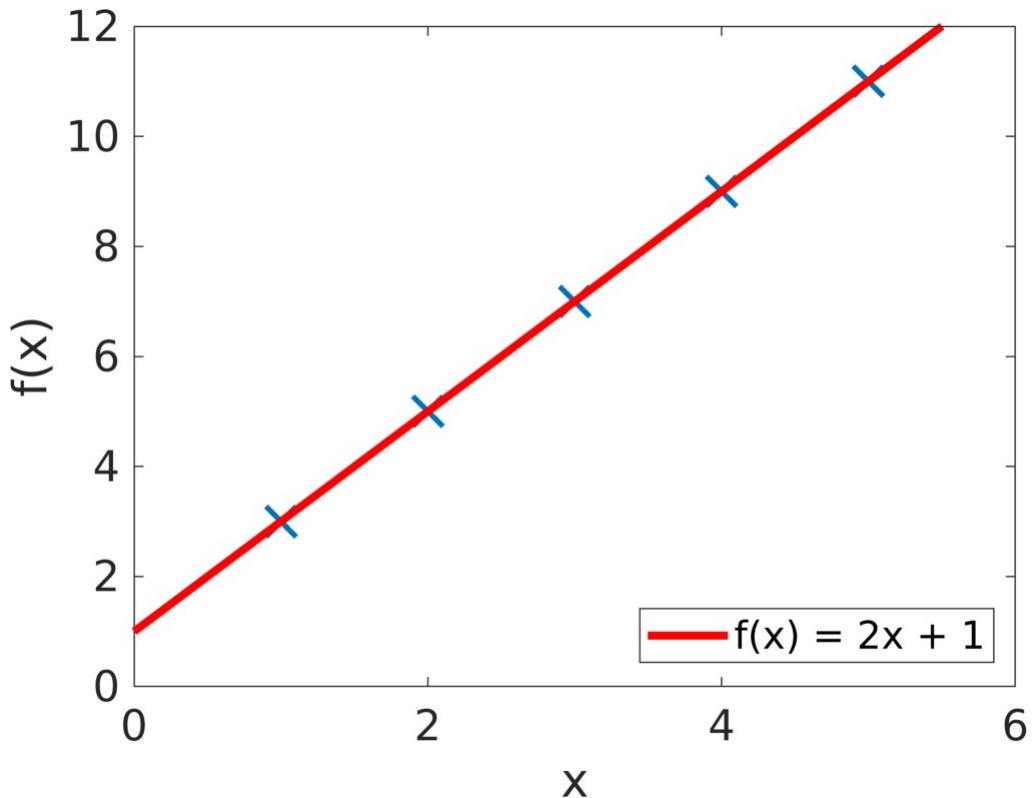
$$x = [1, 2, 3, 4, 5]$$

$$y = [3, 5, 7, 9, 11]$$

$$f(x) = kx + d$$



Data



$$x = [1, 2, 3, 4, 5]$$

$$y = [3, 5, 7, 9, 11]$$

$$f(x) = kx + d$$

$$f(x) = 2x + 1$$

Linear regression with gradient descent: PyTorch

```
## training data
x_train = torch.Tensor([1,2,3,4,5]).float()
y_train = torch.Tensor([3,5,7,9,11]).float()

## model parameters
k = torch.Tensor([0.1]).float()
k.requires_grad=True

d = torch.Tensor([-0.1]).float()
d.requires_grad=True

## model and optimizer
def forward(x):
    return x * k + d

def loss_fcn(x,y):
    y_pred = forward(x)
    return (y_pred-y) * (y_pred-y)

## train
training_epochs = 1000
lr = 0.005
loss_ii = np.zeros(training_epochs)
k_ii = np.zeros(training_epochs)
d_ii = np.zeros(training_epochs)

t = time.time()
for ii in range(training_epochs):
    for x_train_batch, y_train_batch in zip(x_train,y_train):
        loss = loss_fcn(x_train_batch, y_train_batch)
        loss_ii[ii] = loss.item()
        k_ii[ii] = k.data
        d_ii[ii] = d.data
        loss.backward()
        k.data = k.data - lr * k.grad.data
        d.data = d.data - lr * d.grad.data

    # Manually set gradient to zero after update step to prevent gradient accumulation
    k.grad.data.zero_()
    d.grad.data.zero_()
print('epoch: {}, k={:.3}, d={:.3}, loss={:.3}'.format(ii+1,k_ii[ii], d_ii[ii], loss_ii[ii]))

elapsed = time.time() - t
print('Training time: {:.2} s'.format(elapsed))
```



Model and loss function

```
## training data
x_train = torch.Tensor([1,2,3,4,5]).float()
y_train = torch.Tensor([3,5,7,9,11]).float()

## model parameters
k = torch.Tensor([0.1]).float()
k.requires_grad=True

d = torch.Tensor([-0.1]).float()
d.requires_grad=True

## model and optimizer
def forward(x):
    return x * k + d

def loss_fcn(x,y):
    y_pred = forward(x)
    return (y_pred-y) * (y_pred-y)

## train
training_epochs = 1000
lr = 0.005
loss_ii = np.zeros(training_epochs)
k_ii = np.zeros(training_epochs)
d_ii = np.zeros(training_epochs)

t = time.time()
for ii in range(training_epochs):
    for x_train_batch, y_train_batch in zip(x_train,y_train):
        loss = loss_fcn(x_train_batch, y_train_batch)
        loss_ii[ii] = loss.item()
        k_ii[ii] = k.data
        d_ii[ii] = d.data
        loss.backward()
        k.data = k.data - lr * k.grad.data
        d.data = d.data - lr * d.grad.data

    # Manually set gradient to zero after update step to prevent gradient accumulation
    k.grad.data.zero_()
    d.grad.data.zero_()
print('epoch: {}, k={:.3}, d={:.3}, loss={:.3}'.format(ii+1,k_ii[ii], d_ii[ii], loss_ii[ii]))

elapsed = time.time() - t
print('Training time: {:.2} s'.format(elapsed))
```



Training the model

```
## training data
x_train = torch.Tensor([1,2,3,4,5]).float()
y_train = torch.Tensor([3,5,7,9,11]).float()

## model parameters
k = torch.Tensor([0.1]).float()
k.requires_grad=True

d = torch.Tensor([-0.1]).float()
d.requires_grad=True

## model and optimizer
def forward(x):
    return x * k + d

def loss_fcn(x,y):
    y_pred = forward(x)
    return (y_pred-y) * (y_pred-y)

## train
training_epochs = 1000
lr = 0.005
loss_ii = np.zeros(training_epochs)
k_ii = np.zeros(training_epochs)
d_ii = np.zeros(training_epochs)

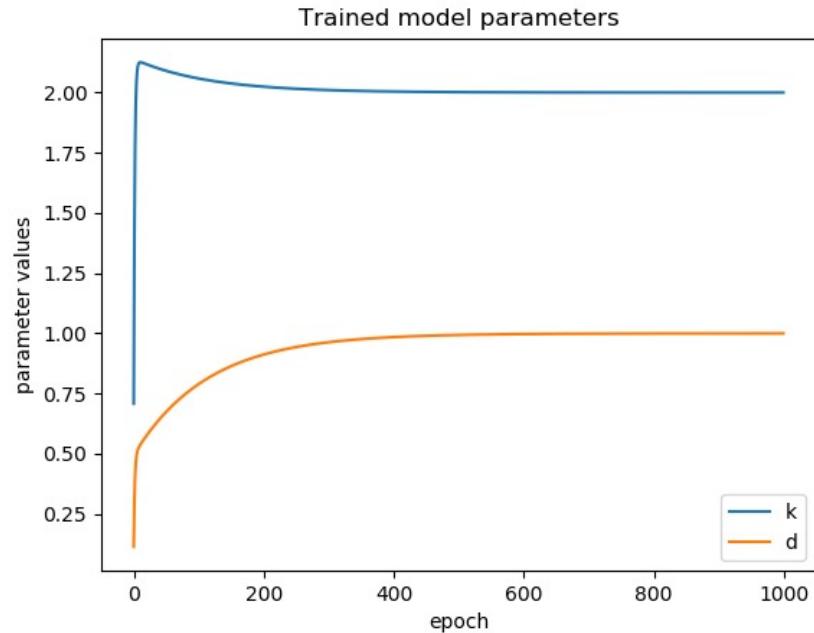
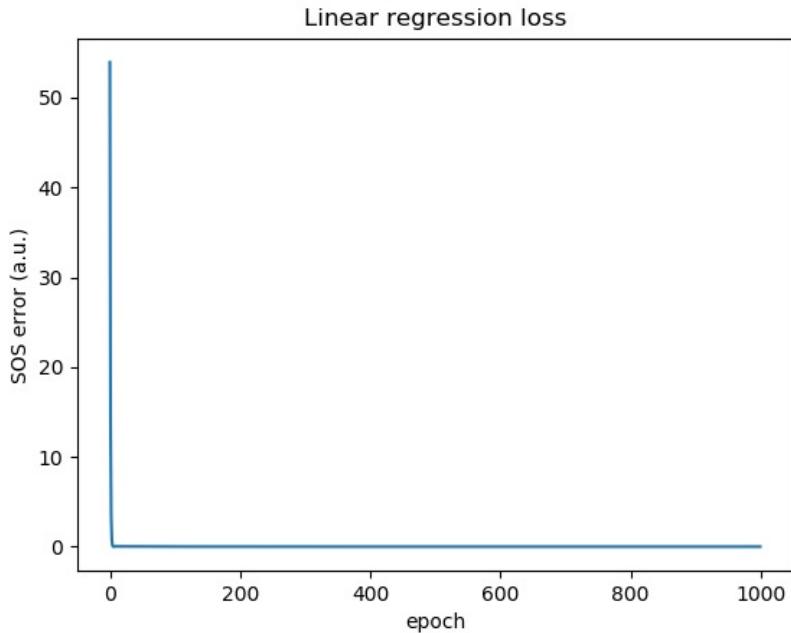
t = time.time()
for ii in range(training_epochs):
    for x_train_batch, y_train_batch in zip(x_train,y_train):
        loss = loss_fcn(x_train_batch, y_train_batch)
        loss_ii[ii] = loss.item()
        k_ii[ii] = k.data
        d_ii[ii] = d.data
        loss.backward()
        k.data = k.data - lr * k.grad.data
        d.data = d.data - lr * d.grad.data

        # Manually set gradient to zero after update step to prevent gradient accumulation
        k.grad.data.zero_()
        d.grad.data.zero_()
    print('epoch: {}, k={:.3}, d={:.3}, loss={:.3}'.format(ii+1,k_ii[ii], d_ii[ii], loss_ii[ii]))

elapsed = time.time() - t
print('Training time: {:.2} s'.format(elapsed))
```



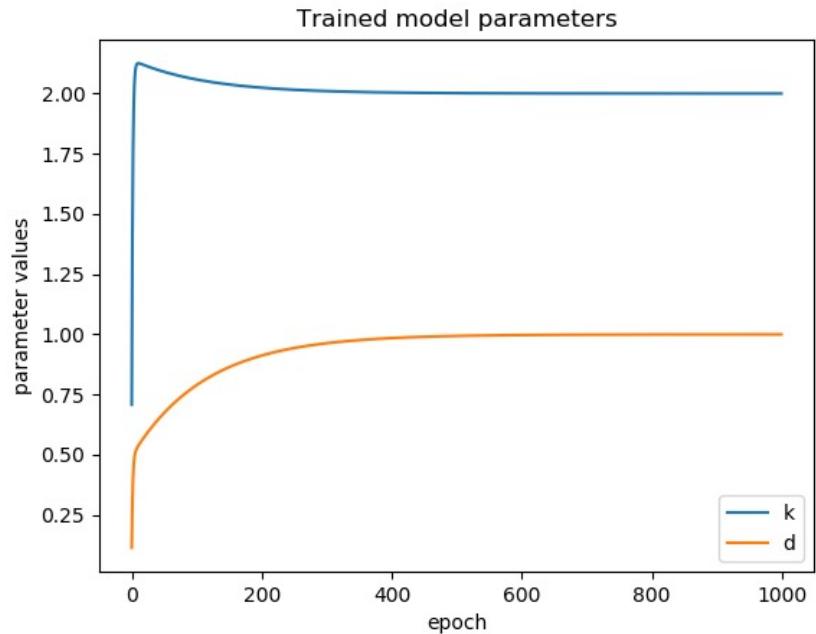
Linear regression with gradient descent: PyTorch



Linear regression in PyTorch

```
epoch: 993, k=2.0, d=1.0, loss=1.11e-09
epoch: 994, k=2.0, d=1.0, loss=1.11e-09
epoch: 995, k=2.0, d=1.0, loss=9.9e-10
epoch: 996, k=2.0, d=1.0, loss=9.9e-10
epoch: 997, k=2.0, d=1.0, loss=9.9e-10
epoch: 998, k=2.0, d=1.0, loss=9.9e-10
epoch: 999, k=2.0, d=1.0, loss=9.31e-10
epoch: 1000, k=2.0, d=1.0, loss=9.9e-10
Training time: 1.4 s
```

$$f(x) = 2x + 1$$



Example 2: Classification of brain tissue from DTI data

High level neural network modules

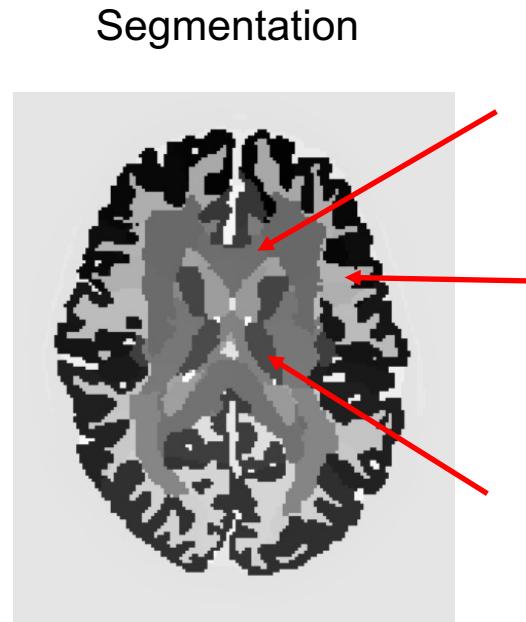
Classification of brain tissue from HCP DTI data



Genu of corpus callusum
(highly aligned WM)

Subcortical WM

Thalamus (GM)



Classification of brain tissue from HCP DTI data

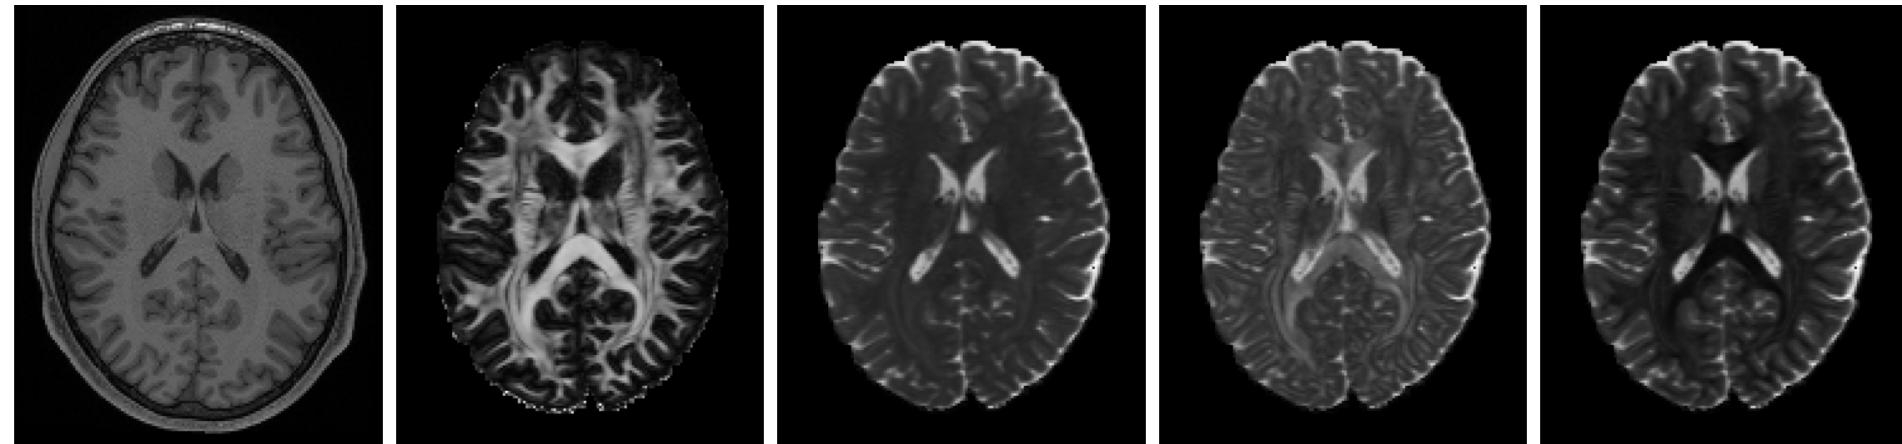
T1w

FA

MD $(\lambda_1 + \lambda_2 + \lambda_3)/3$

AD (λ_1)

RD $(\lambda_2 + \lambda_3)/2$



fMRI

<https://www.humanconnectome.org/>

Structure and inspection of the data

$$\{(x_1, y_1), \dots (x_N, y_N)\}$$

$$x_i = [T1w_i, FA_i, MD_i, AD_i, RD_i]$$

T1w (a.u.)	FA (-)	MD $(\frac{\mu m^2}{ms})$	AD $(\frac{\mu m^2}{ms})$	RD $(\frac{\mu m^2}{ms})$	Class	Class label
898	0.22	1.066592	1.33	0.94	Thalamus	1
1007	0.68	0.39	0.72	0.22	CC	2
867	0.38	0.58	0.82	0.45	Cortical WM	3
...

Normalization!

Training, validation and test set

$$\{(x_1, y_1), \dots (x_N, y_N)\}$$

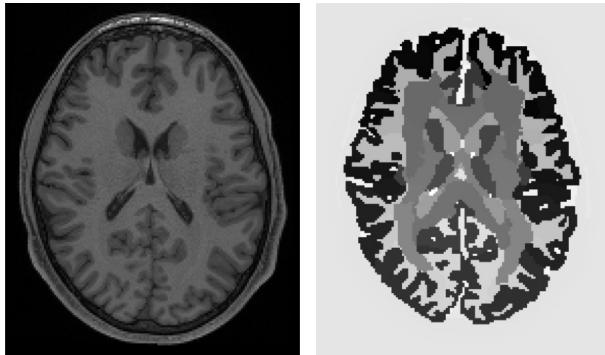
1. Training set: Train the model
2. Validation set: Use as performance criterion to tune training, for hyperparameter selection,...
3. Test set: Don't use during training at all!

Training, validation and test set

1. Training set: 80% randomized samples from 3 subjects
2. Validation set: 20% randomized samples from 3 subjects
3. Test set: Randomized samples from 1 different subject

$$\{(x_1, y_1), \dots (x_N, y_N)\} \quad N = 23700 \quad N_{train} = 14044 \\ N_{val} = 3511 \quad N_{test} = 6145$$

Check balancing of classes



$$N_{thal} = 11662$$

$$N_{cc} = 4582$$

$$N_{subcort-wm} = 7455$$

- Create data duplicates
- Weight samples during training

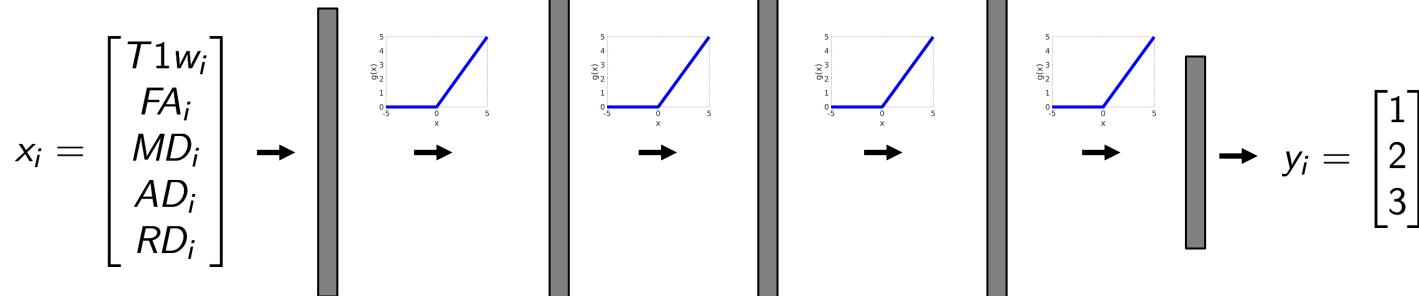
Let's build our first neural network

Multi layer perceptron

- 1 input layer
- 3 hidden layers with 100 elements
- 1 output layer
- Non-linearity: ReLu
- Softmax output

Model Architecture

$$y_i(\mathbf{x}) = g(\mathbf{W}_i \mathbf{x} + \mathbf{b}_i)$$



$$(5*100)+100 = 600$$

$$(100*3)+3 = 303$$

Total number of parameters: 31203

$$(100*100)+100 = 10100$$

PyTorch implementation

```
%% Define model
nElements = 100
nLayers = 3
model_name = 'dti_FC'
model = torch.nn.Sequential(
    torch.nn.Linear(nFeatures, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nClasses, bias=True),
)
print(model)

%%choose optimizer and loss function
training_epochs = 250
lr = 0.001
batch_size = 1024
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=lr)

%%Create minibatch data loading for training and validation
dataLoader_train = data_utils.TensorDataset(x_train, y_train)
dataLoader_train = data_utils.DataLoader(dataLoader_train, batch_size=batch_size, shuffle=False,num_workers=4)

%% Train model
loss_train = np.zeros(training_epochs)
acc_train = np.zeros(training_epochs)
loss_val = np.zeros(training_epochs)
acc_val = np.zeros(training_epochs)

for epoch in range(training_epochs):
    for local_batch, local_labels in dataLoader_train:
        # feedforward - backpropagation
        optimizer.zero_grad()
        out = model(local_batch)
        loss = criterion(out, local_labels)
        loss.backward()
        optimizer.step()
        loss_train[epoch] = loss.item()
        # Training data accuracy
        [dummy, predicted] = torch.max(out.data, 1)
        acc_train[epoch] = (torch.sum(local_labels==predicted).numpy() / np.size(local_labels.numpy(),0))

    # Validation
    out_val = model(x_val)
    loss = criterion(out_val, y_val)
    loss_val[epoch] = loss.item()
    [dummy, predicted_val] = torch.max(out_val.data, 1)
    acc_val[epoch] = ( torch.sum(y_val==predicted_val).numpy() / setsize_val)

print ('Epoch {}/{} train loss: {:.3}, train acc: {:.3}, val loss: {:.3}, val acc: {:.3}'.format(epoch+1, training_epochs, loss_train[epoch], acc_train[epoch], loss_val[epoch], acc_val[epoch]))
```



Set up optimizer and data support functions

```
#% Define model
nElements = 100
nLayers = 3
model_name = 'dti_FC'
model = torch.nn.Sequential(
    torch.nn.Linear(nFeatures, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nElements, bias=True),
    torch.nn.ReLU(),
    torch.nn.Linear(nElements, nClasses, bias=True),
)
print(model)

#%choose optimizer and loss function
training_epochs = 250
lr = 0.001
batch_size = 1024
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=lr)

#%Create minibatch data loading for training and validation
dataLoader_train = data_utils.TensorDataset(x_train, y_train)
dataLoader_train = data_utils.DataLoader(dataLoader_train, batch_size=batch_size, shuffle=False,num_workers=4)

#% Train model
loss_train = np.zeros(training_epochs)
acc_train = np.zeros(training_epochs)
loss_val = np.zeros(training_epochs)
acc_val = np.zeros(training_epochs)

for epoch in range(training_epochs):
    for local_batch, local_labels in dataLoader_train:
        # feedforward - backpropagation
        optimizer.zero_grad()
        out = model(local_batch)
        loss = criterion(out, local_labels)
        loss.backward()
        optimizer.step()
        loss_train[epoch] = loss.item()
        # Training data accuracy
        [dummy, predicted] = torch.max(out.data, 1)
        acc_train[epoch] = (torch.sum(local_labels==predicted).numpy() / np.size(local_labels.numpy(),0))

    # Validation
    out_val = model(x_val)
    loss = criterion(out_val, y_val)
    loss_val[epoch] = loss.item()
    [dummy, predicted_val] = torch.max(out_val.data, 1)
    acc_val[epoch] = ( torch.sum(y_val==predicted_val).numpy() / setsize_val)

print ('Epoch {}/{} train loss: {:.3}, train acc: {:.3}, val loss: {:.3}, val acc: {:.3}'.format(epoch+1, training_epochs, loss_train[epoch], acc_train[epoch], loss_val[epoch], acc_val[epoch]))
```



Kingma and Ba, ICLR 2015

Let's train our network

```
%>% Define model
nElements = 100
nLayers = 3
model_name = 'dti_FC'
model = torch.nn.Sequential(
  torch.nn.Linear(nFeatures, nElements, bias=True),
  torch.nn.ReLU(),
  torch.nn.Linear(nElements, nElements, bias=True),
  torch.nn.ReLU(),
  torch.nn.Linear(nElements, nElements, bias=True),
  torch.nn.ReLU(),
  torch.nn.Linear(nElements, nElements, bias=True),
  torch.nn.ReLU(),
  torch.nn.Linear(nElements, nClasses, bias=True),
)
print(model)

#%choose optimizer and loss function
training_epochs = 250
lr = 0.001
batch_size = 1024
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=lr)

#%Create minibatch data loading for training and validation
dataLoader_train = data_utils.TensorDataset(x_train, y_train)
dataLoader_train = data_utils.DataLoader(dataLoader_train, batch_size=batch_size, shuffle=False,num_workers=4)

#% Train model
loss_train = np.zeros(training_epochs)
acc_train = np.zeros(training_epochs)
loss_val = np.zeros(training_epochs)
acc_val = np.zeros(training_epochs)

for epoch in range(training_epochs):
  for local_batch, local_labels in dataLoader_train:
    # feedforward - backpropagation
    optimizer.zero_grad()
    out = model(local_batch)
    loss = criterion(out, local_labels)
    loss.backward()
    optimizer.step()
    loss_train[epoch] = loss.item()
    # Training data accuracy
    [dummy, predicted] = torch.max(out.data, 1)
    acc_train[epoch] = (torch.sum(local_labels==predicted).numpy() / np.size(local_labels.numpy(),0))

    # Validation
    out_val = model(x_val)
    loss = criterion(out_val, y_val)
    loss_val[epoch] = loss.item()
    [dummy, predicted_val] = torch.max(out_val.data, 1)
    acc_val[epoch] = ( torch.sum(y_val==predicted_val).numpy() / setsize_val)

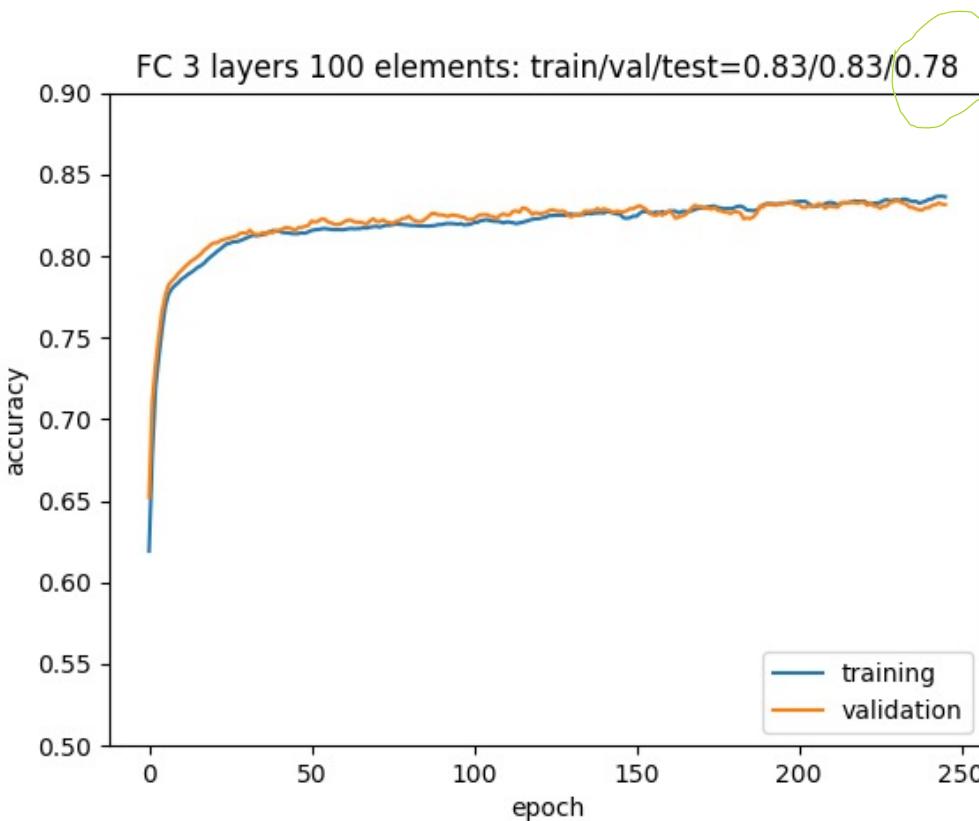
  print ('Epoch {}/{} train loss: {:.3}, train acc: {:.3}, val loss: {:.3}, val acc: {:.3}'.format(epoch+1, training_epochs, loss_train[epoch], acc_train[epoch], loss_val[epoch], acc_val[epoch]))
```



Kingma and Ba, ICLR 2015

Results

test set里面有其他病人，模型的泛化能力不够好

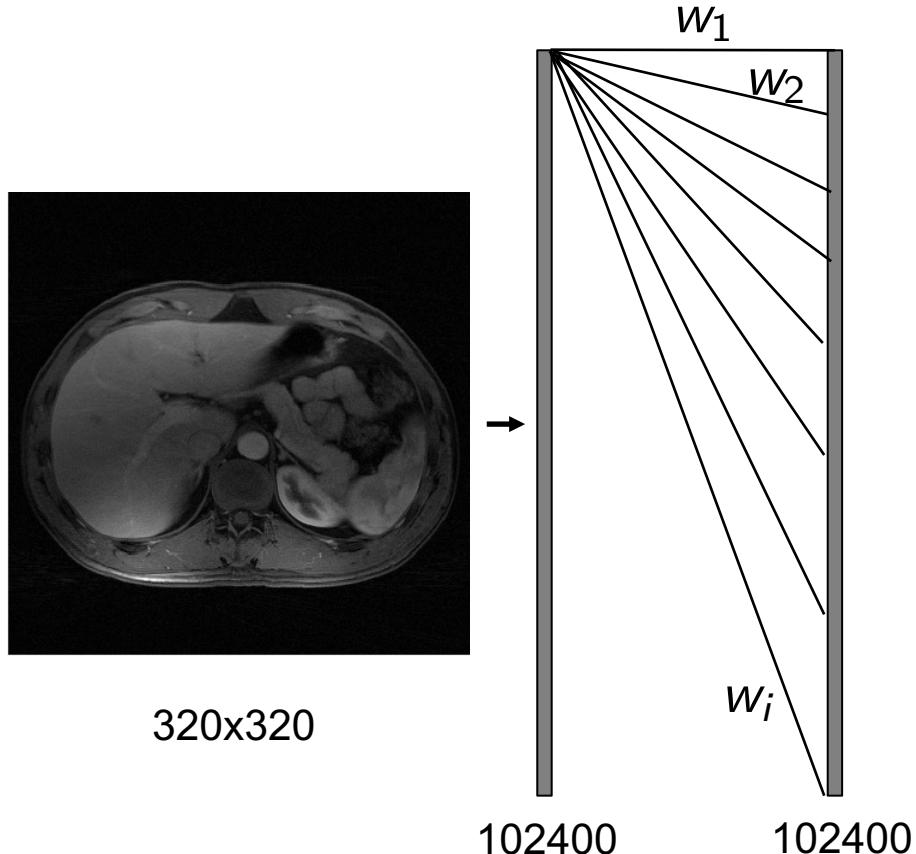


- Higher model complexity?
- Inspect misclassified samples
- Trends for different label categories?
- Experiment with features
- Don't use test performance to adjust model architecture!

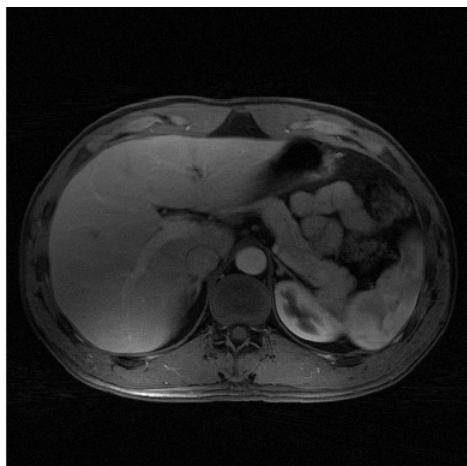
Example 3: Classification of image quality of
accelerated reconstructions

Convolutional Neural Networks (CNNs),
model complexity

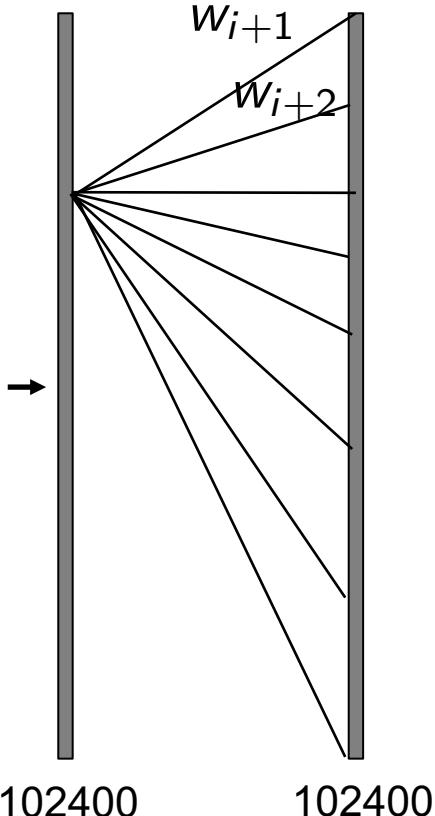
Fully connected Neural Networks



Fully connected Neural Networks



320x320

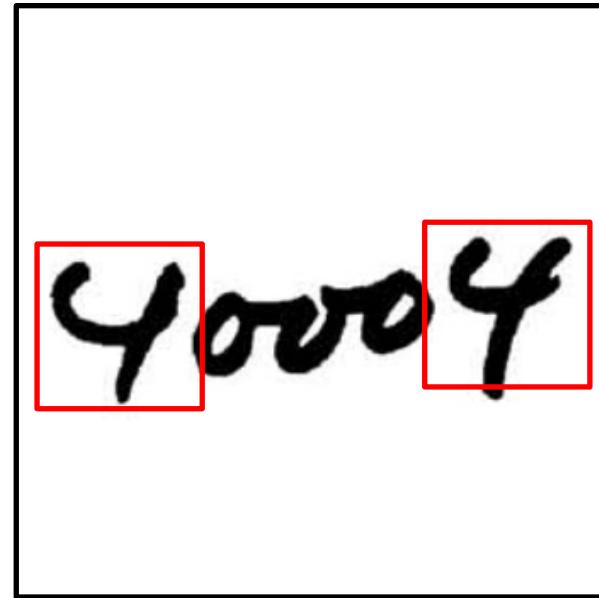


Number of parameters:
 $102400 * 102400 \approx 1.05 * 10^{10}$

**More efficient use of
parameters!**

Convolutional neural networks

40004 75216
14199-2087 23505
96203 14310
44151 .05153



Local connectivity
Share parameters

Convolutional neural networks

40004 75216
14199-2087 23505
96203 14310
44151 .05153

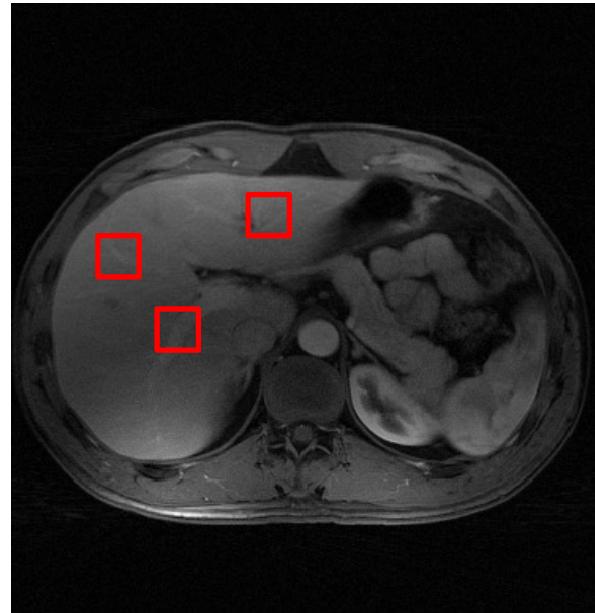
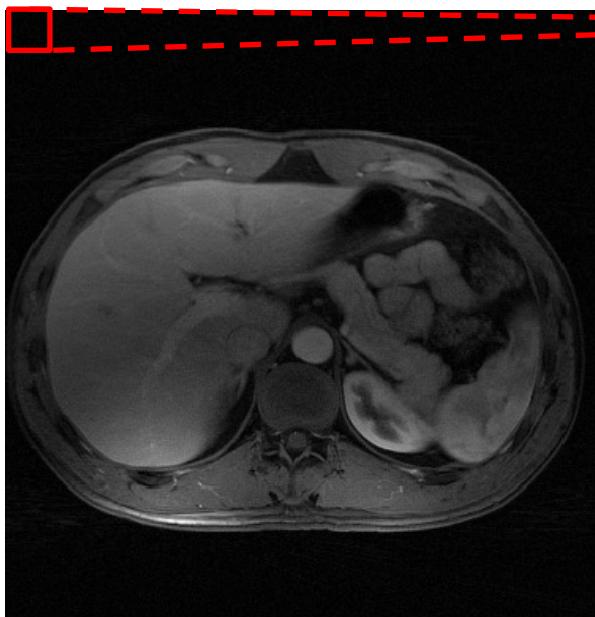


Figure 1: Examples of original zipcodes from the testing set.

Local connectivity
Share parameters

Convolutional layers

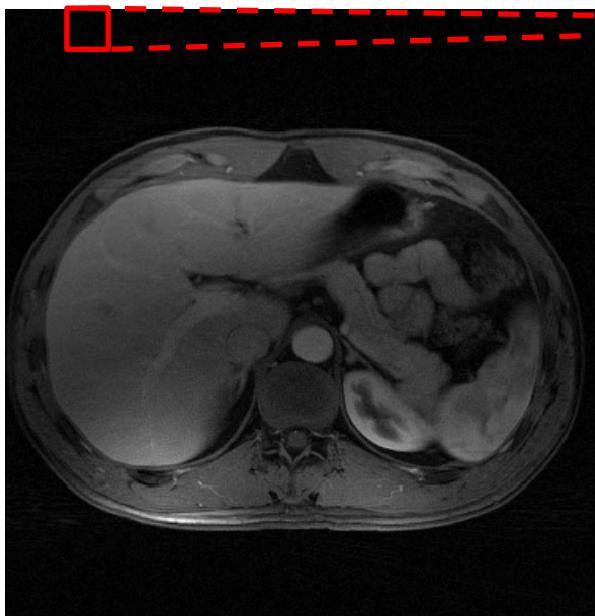
$$w^T x + b$$



320x320 image
3x3 filter w

Convolutional layers

$$w^T x + b$$

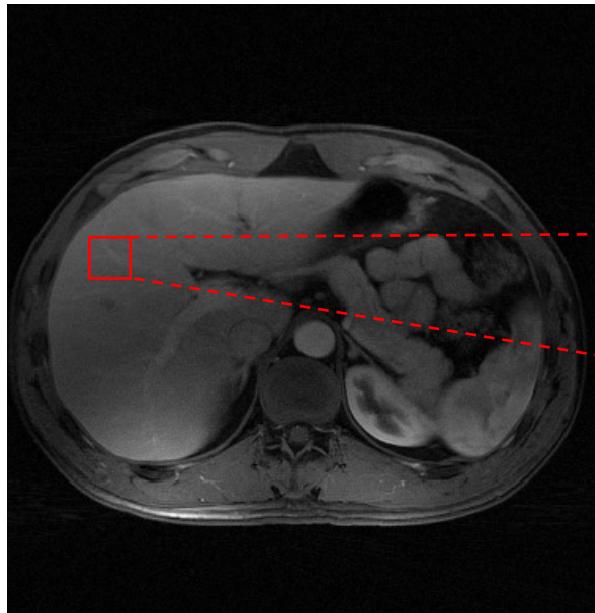


320x320 image
3x3 filter w



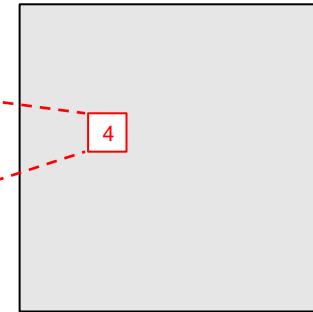
Model parameters: $3 \times 3 + 1 = 10$

Pooling layers

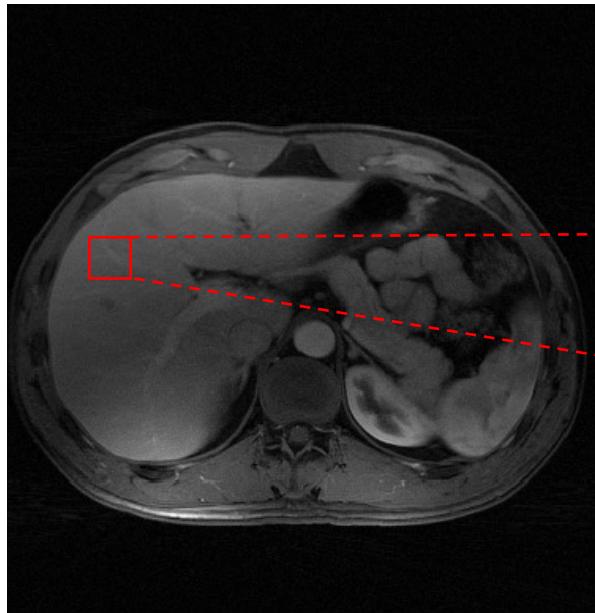


2x2 max pooling

1	2
3	4



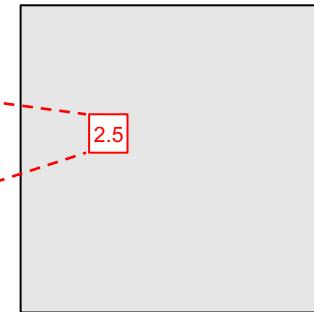
Pooling layers



320x320 image

2x2 avg pooling

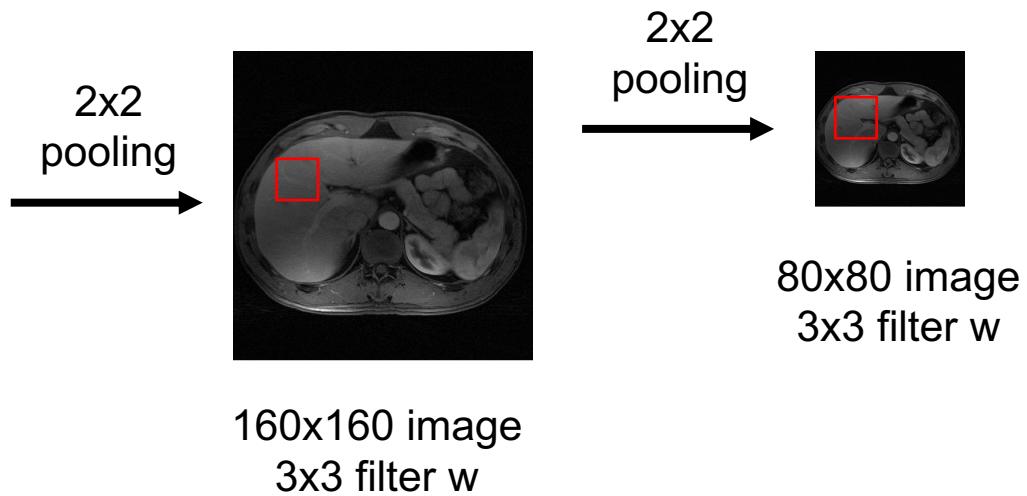
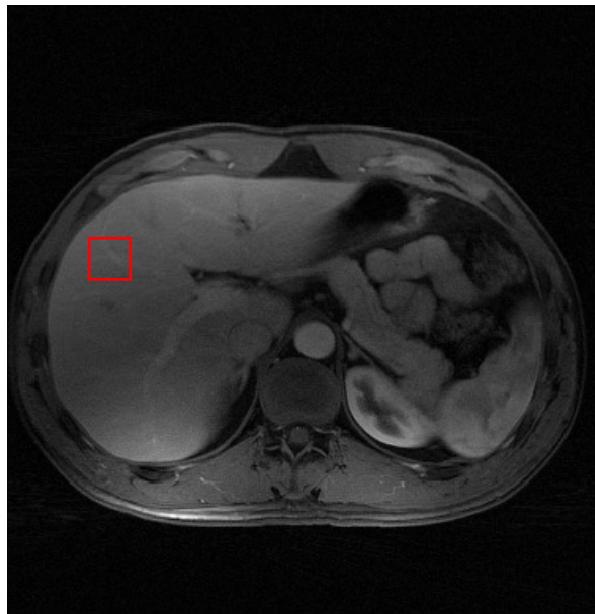
1	2
3	4



160x160 image

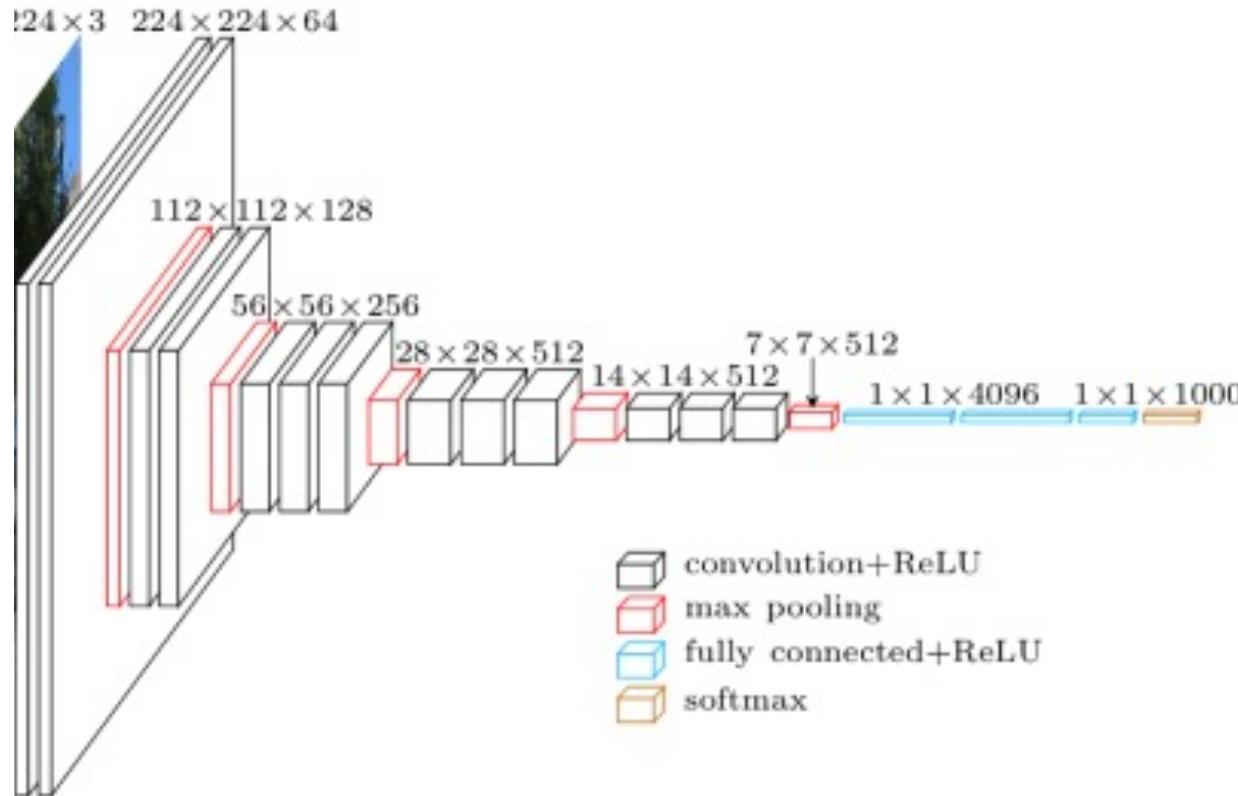
No model parameters

Pooling layers



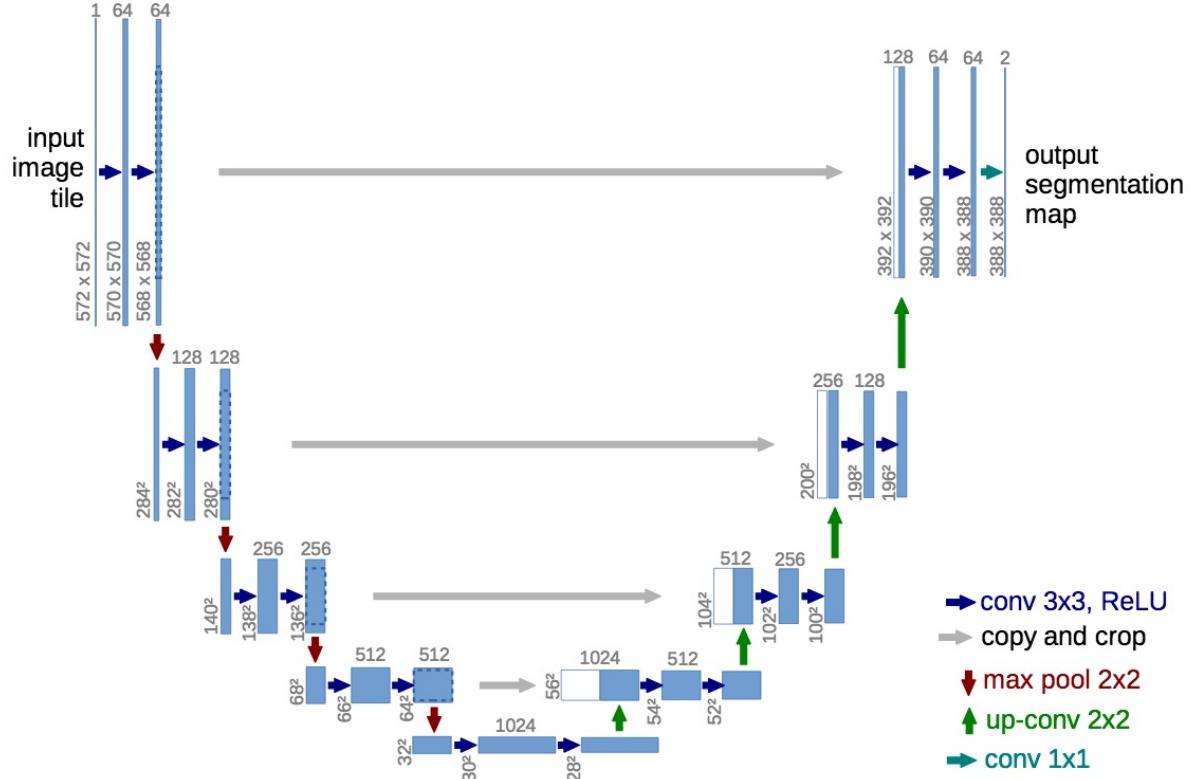
- Increases receptive field
- Multi scale image analysis
- Compression of information

Example architecture: VGG



Popular for image categorization

Example architecture: U-Net



Popular for image
processing
and segmentation

Example 3: Classification of image quality of accelerated reconstructions

Fully sampled vs 4 times PI-CS accelerated

CS:compress sensing

Fully sampled reference



PI-CS (TGV) R=4



Fully sampled vs 4 times PI-CS accelerated

Fully sampled reference



PI-CS (TGV) R=4



Fully sampled vs 4 times PI-CS accelerated

Fully sampled reference

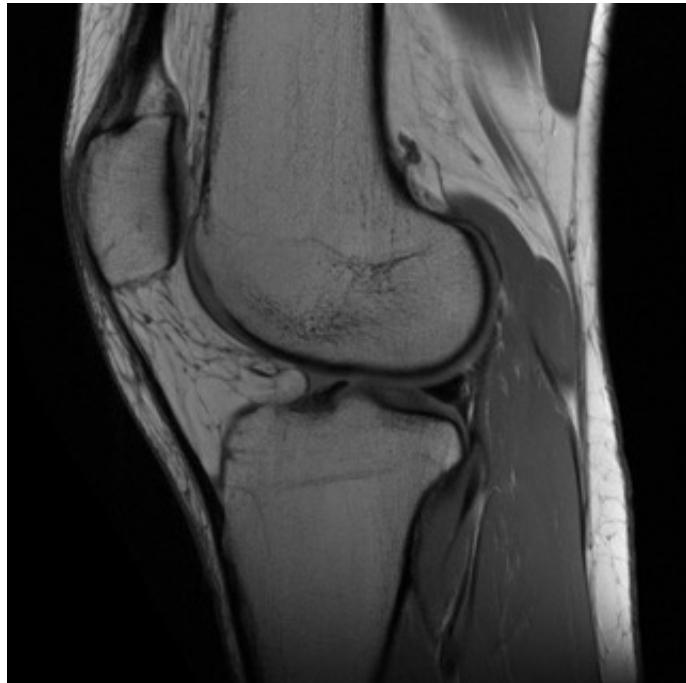


PI-CS (TGV) R=4

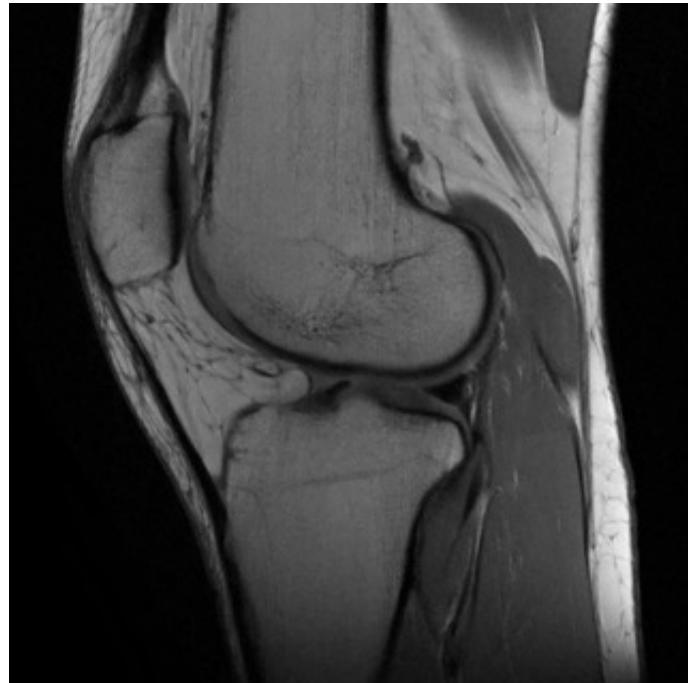


Fully sampled vs 4 times PI-CS accelerated

Fully sampled reference



PI-CS (TGV) R=4



Training, validation and test set

1. Training set: 2/3 randomized samples
2. Validation set: 1/3 randomized samples
3. No separate test set: Experiment is left to you :-)

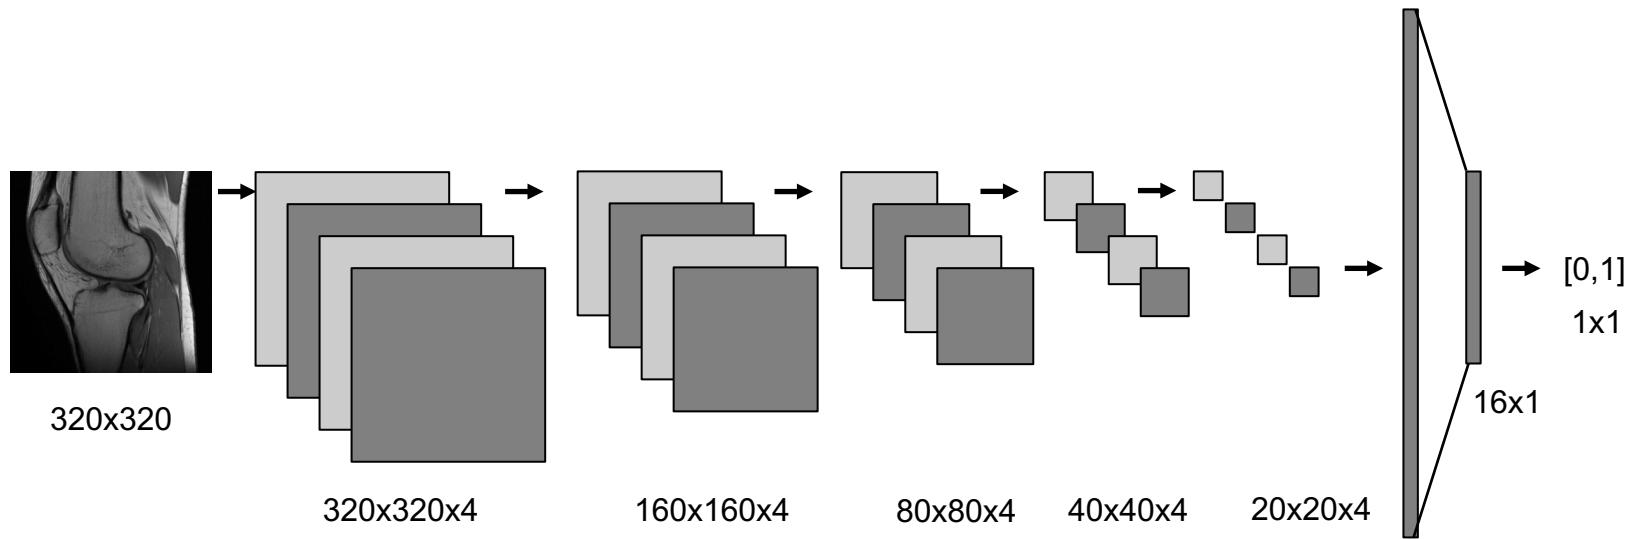
$$N = 520$$

$$\{(x_1, y_1), \dots (x_N, y_N)\} \quad N_{train} = 390 \quad N_{ref} = 260$$

$$N_{val} = 130 \quad N_{tgv} = 260$$

total generalize validation?

Our first convolutional neural network (CNN)



4 $3 \times 3 \times 4$ convolutional layers, Relu activation

2D MaxPooling

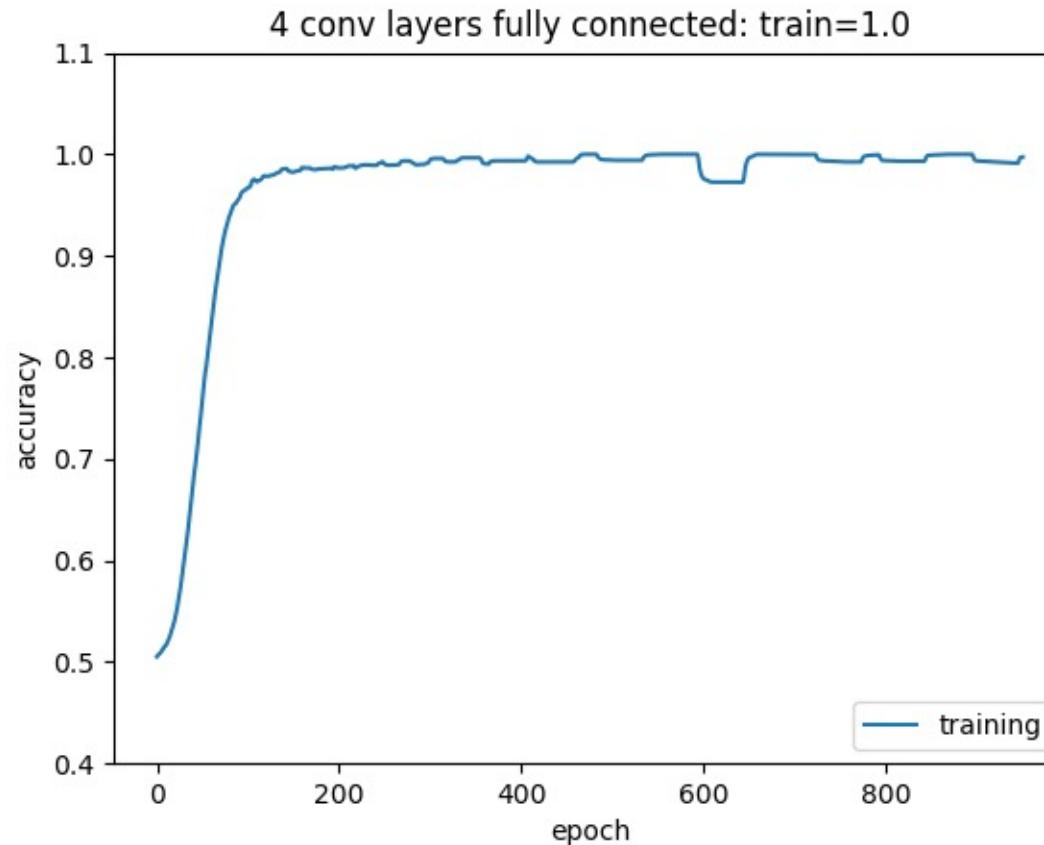
1 fully connected layer with 16 elements

Sigmoid Output

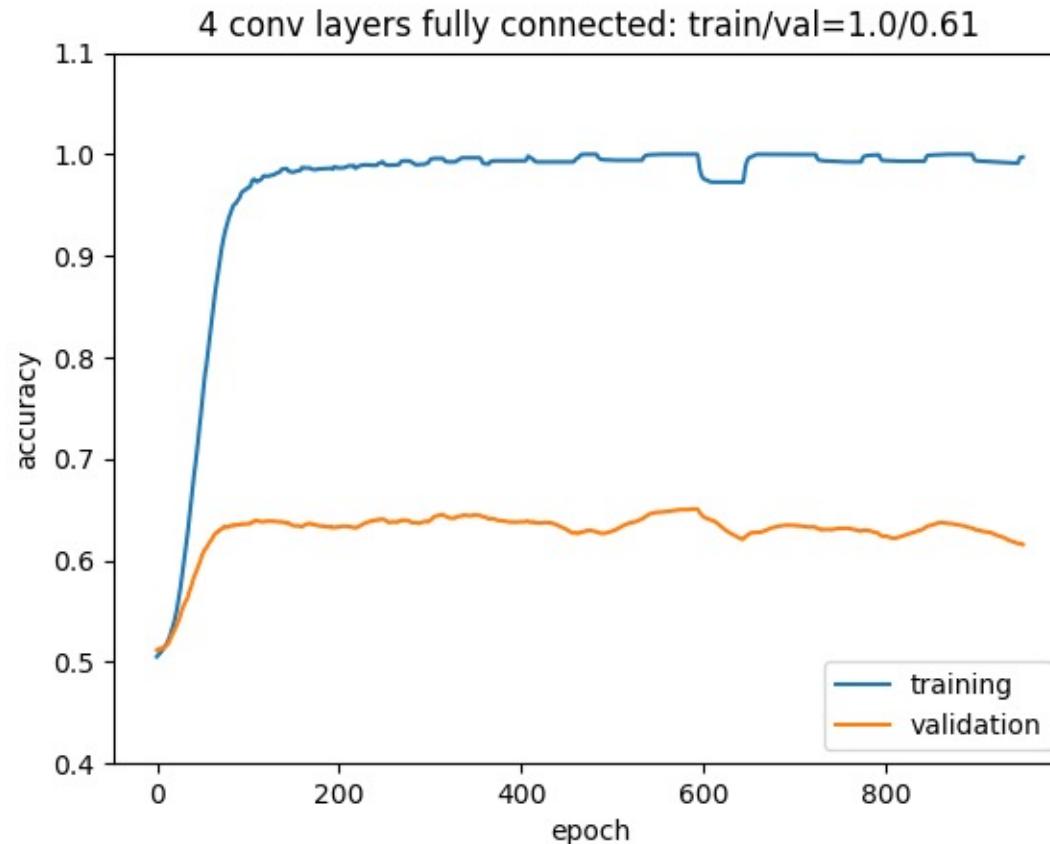
1600x1

Total number of parameters: 26117

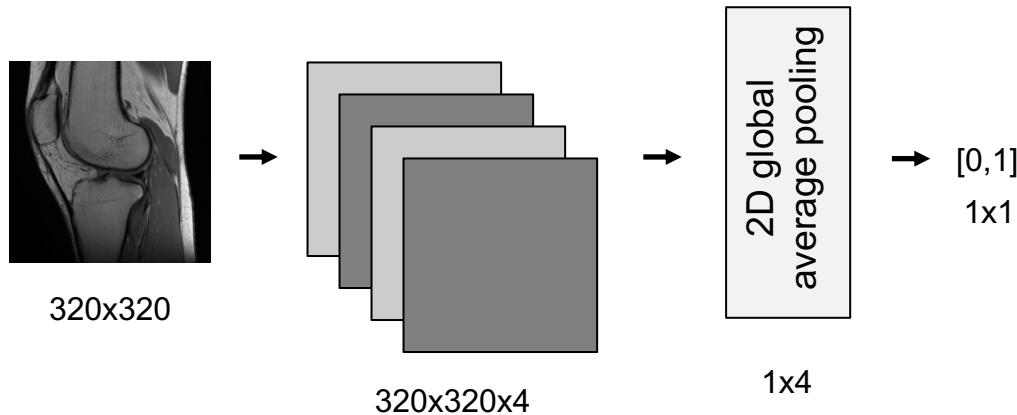
Results



Results



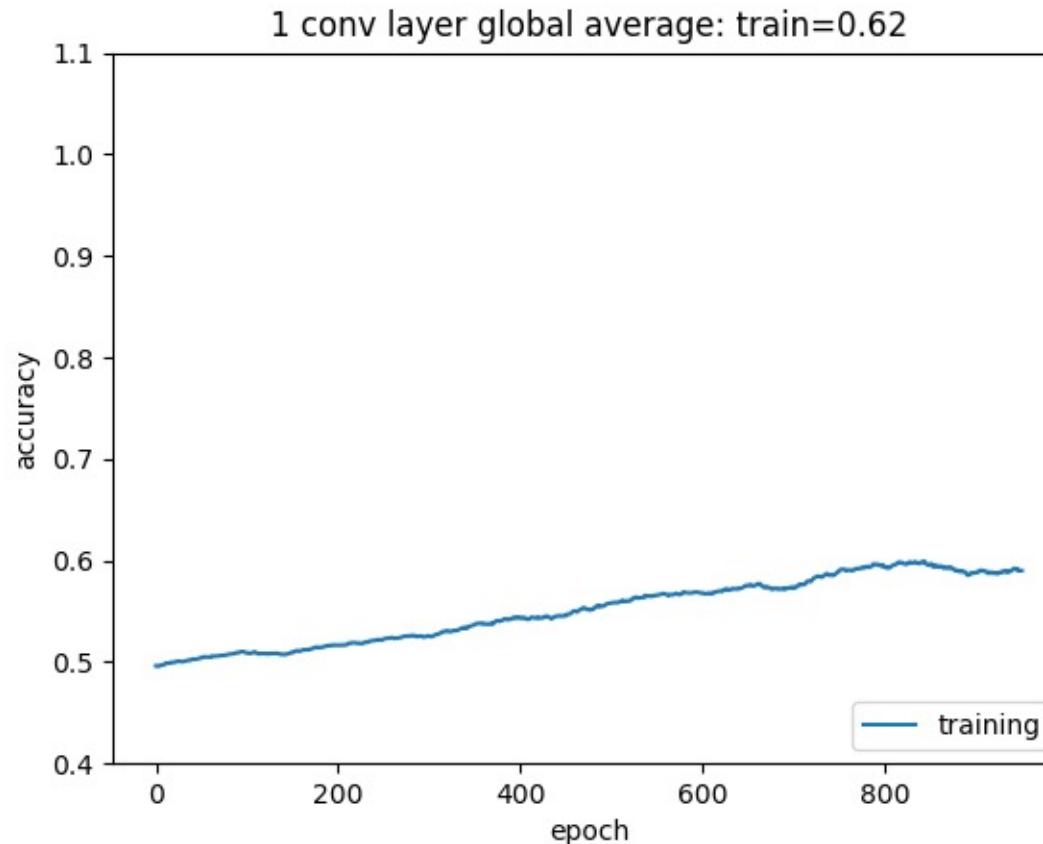
Reduce complexity of our model



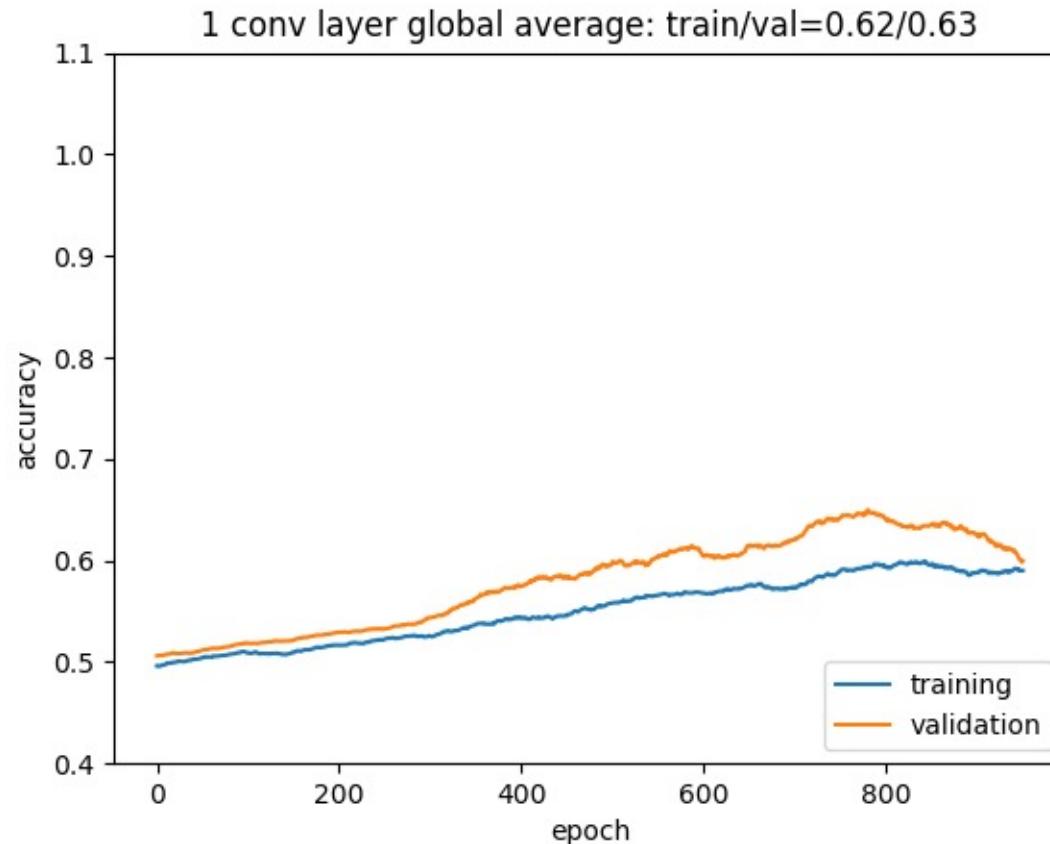
1 $3 \times 3 \times 4$ convolutional layer, Relu activation
2D Global Average Pooling
Sigmoid Output

Total number of parameters: 45

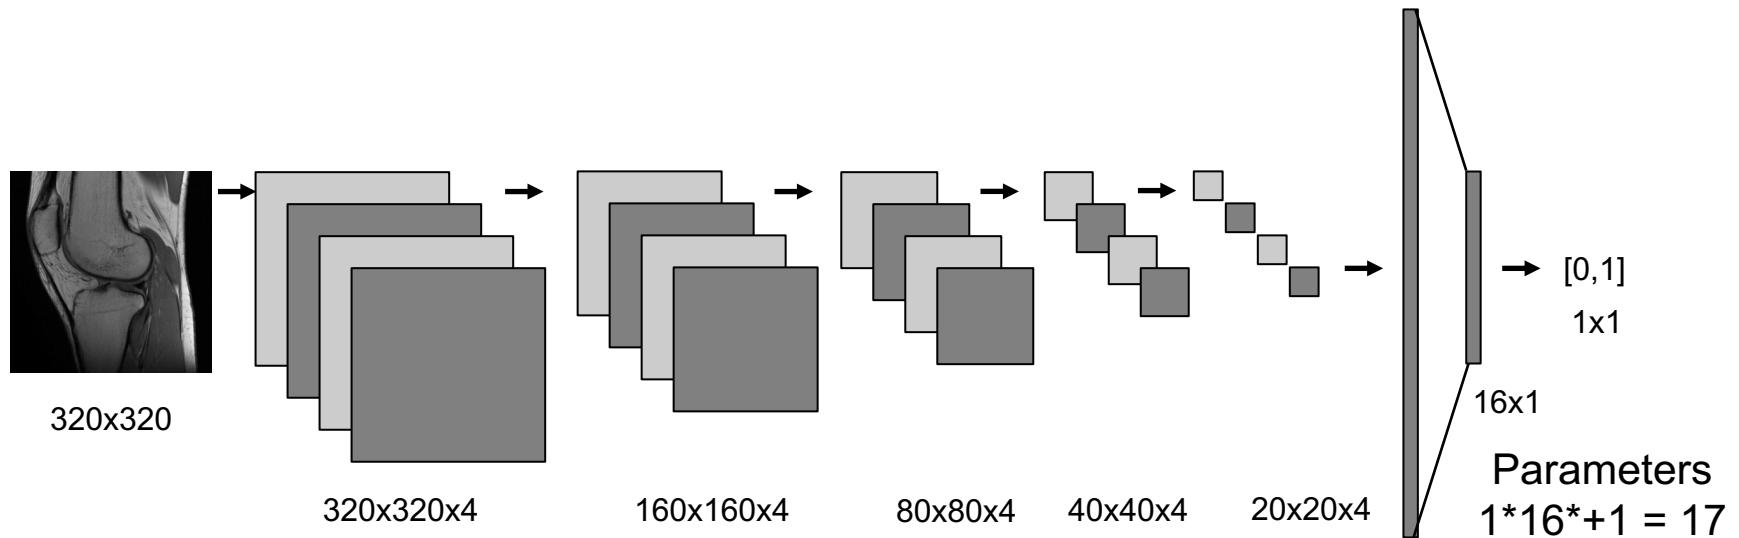
Results



Results



A closer look at the parameter distribution of the first model



Parameters
 $(3 \times 3 + 1) \times 4 = 40$

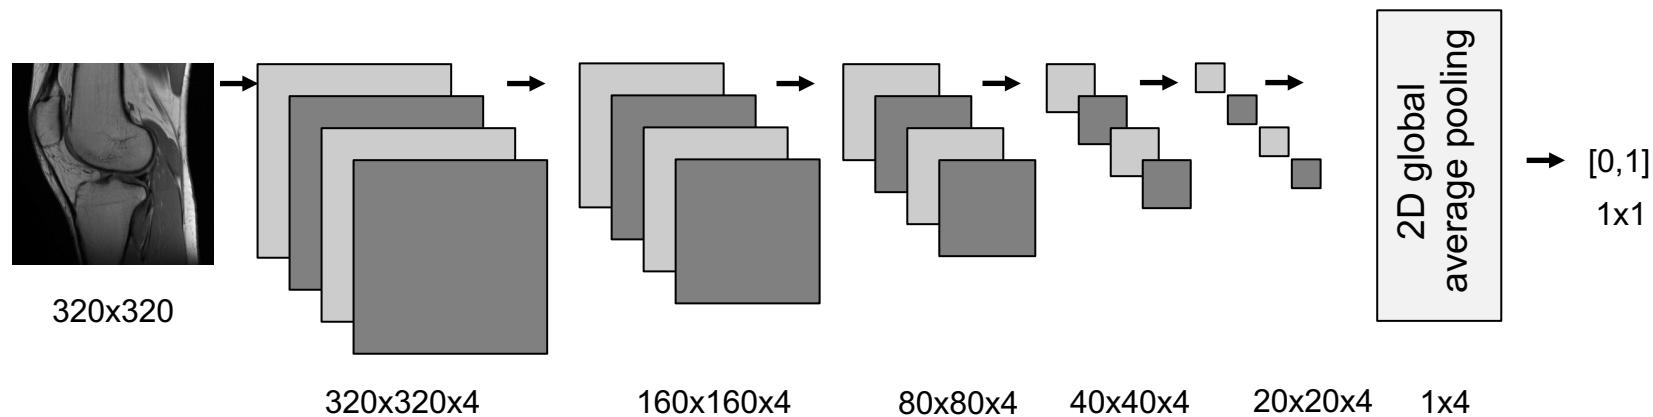
Parameters
 $(3 \times 3 \times 4 + 1) \times 4 = 148$

1600x1

Parameters
 $20 \times 20 \times 4 \times 16 + 16 = 25616$

98% of parameters in fully connected layer!

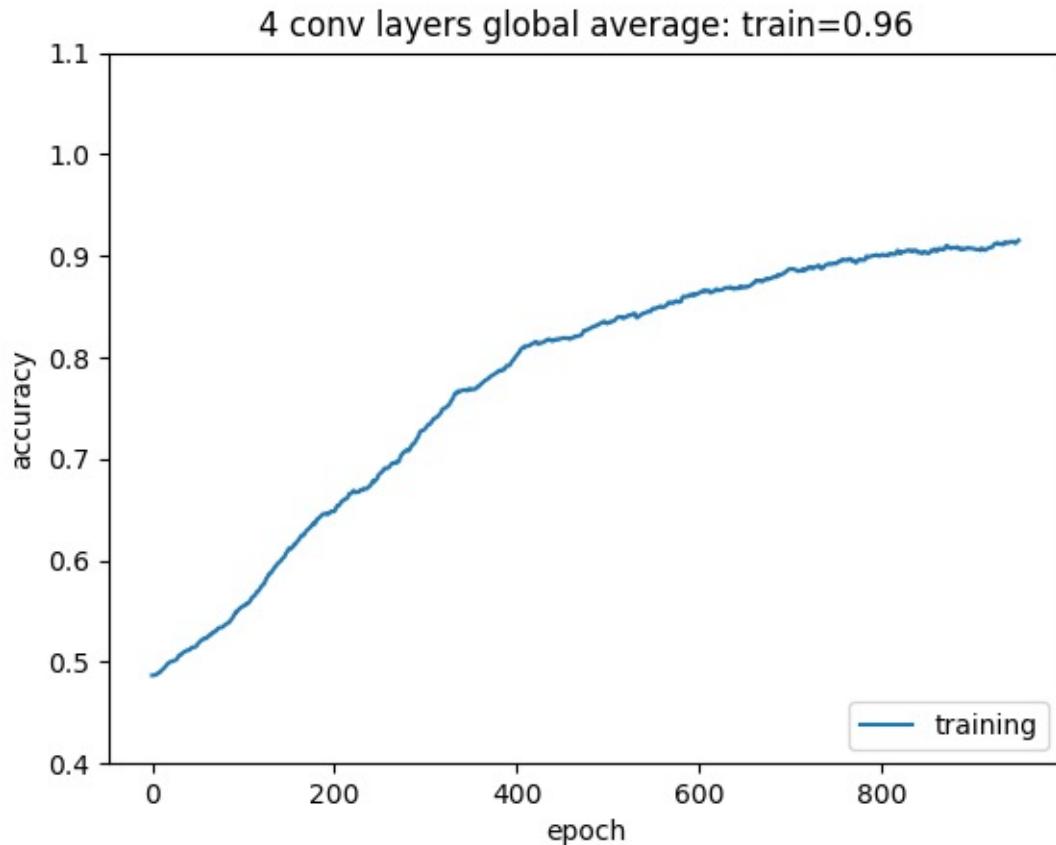
Replace FC layer with global average pooling



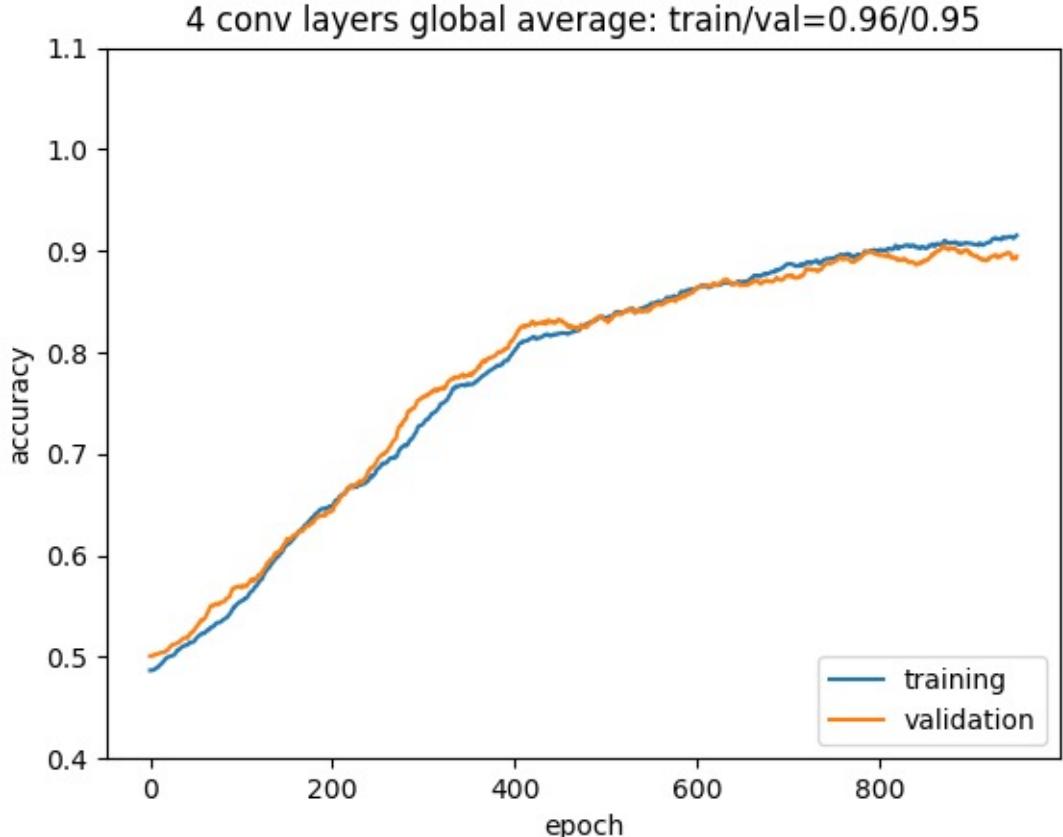
4 3x3x4 convolutional layers, Relu activation
2D MaxPooling
2D Global Average Pooling

Total number of parameters: 489

Results



Results



independent test set

- Use separate test set
- Inspect misclassified samples
- Number of conv layers, filter size,...
- Use higher model complexity combined with regularization (e.g. dropout)

Srivastava et al., JMLR (2014)

Summary

- Introduction to DL/ML/AI
- Examples for linear regression, MLPs, CNNs
- High level view on generalization
- Model complexity, overfitting vs. underfitting
- Model parameters: Convolutional vs fully connected layers